

- 如果在一个周期里面,剩余运行时间用光了,可以尝试那补偿5ms的时间,即总的运行时间减少了5ms,而剩余运行时间增加了5ms.
- 随着时间的流逝,剩余运行时间逐渐减少到0甚至为负值,如果在检测过程中,检测到了剩余运行时间已经使用完毕,那么系统就会额外的补偿给剩余运行时间数值为 5-runtime_remaining(unit:ms).
- 在每次pick task的时候都会检测是否可以throttle,如果可以,则强制将enqueue的task dequeue,并有一个period timer(100ms)定时检测是否有rq throttle了,如果有则cfs调度算法重新对task进行调度操作.

下面是它的初始化:

```

• /*执行slack_timer的回调函数*/
• static enum hrtimer_restart sched_cfs_slack_timer(struct hrtimer *timer)
• {
•     struct cfs_bandwidth *cfs_b =
•         container_of(timer, struct cfs_bandwidth, slack_timer);
•
•     do_sched_cfs_slack_timer(cfs_b);
•
•     return HRTIMER_NORESTART;
• }
• /*running period timer*/
• static enum hrtimer_restart sched_cfs_period_timer(struct hrtimer *timer)
• {
•     struct cfs_bandwidth *cfs_b =
•         container_of(timer, struct cfs_bandwidth, period_timer);
•     int overrun;
•     int idle = 0;
•
•     raw_spin_lock(&cfs_b->lock);
•     for (;;) {
•         overrun = hrtimer_forward_now(timer, cfs_b->period);
•         if (!overrun)
•             break;
•
•         idle = do_sched_cfs_period_timer(cfs_b, overrun);
•     }
•     if (idle)
•         cfs_b->period_active = 0;
•     raw_spin_unlock(&cfs_b->lock);
•
•     return idle ? HRTIMER_NORESTART : HRTIMER_RESTART;
• }
• /*
•  * default period for cfs group bandwidth.
•  * default: 0.1s, units: nanoseconds
•  */
• static inline u64 default_cfs_period(void)
• {
•     return 100000000ULL;
• }
• /*cfs bandwidth的初始化*/
• void init_cfs_bandwidth(struct cfs_bandwidth *cfs_b)

```

```

• {
•     raw_spin_lock_init(&cfs_b->lock);
•     cfs_b->runtime = 0;
•     cfs_b->quota = RUNTIME_INF;
•     cfs_b->period = ns_to_ktime(default_cfs_period());
•
•     INIT_LIST_HEAD(&cfs_b->throttled_cfs_rq);
•     /*周期性处理cfs bandwidth上的task_group*/
•     hrtimer_init(&cfs_b->period_timer, CLOCK_MONOTONIC,
HRTIMER_MODE_ABS_PINNED);
•     cfs_b->period_timer.function = sched_cfs_period_timer;
•     hrtimer_init(&cfs_b->slack_timer, CLOCK_MONOTONIC, HRTIMER_MODE_REL);
•     cfs_b->slack_timer.function = sched_cfs_slack_timer;
• }
•
• static void init_cfs_rq_runtime(struct cfs_rq *cfs_rq)
• {
•     cfs_rq->runtime_enabled = 0;
•     INIT_LIST_HEAD(&cfs_rq->throttled_list);
• }
•
• void start_cfs_bandwidth(struct cfs_bandwidth *cfs_b)
• {
•     lockdep_assert_held(&cfs_b->lock);
•
•     if (!cfs_b->period_active) {
•         cfs_b->period_active = 1;
•         hrtimer_forward_now(&cfs_b->period_timer, cfs_b->period);
•         hrtimer_start_expires(&cfs_b->period_timer,
HRTIMER_MODE_ABS_PINNED);
•     }
• }
• }

```

2. 因为一个task_group是在percpu上都创建了一个cfs_rq，所以cfs_b->quota的值是这些percpu cfs_rq中的进程共享的，每个percpu cfs_rq在运行时需要向tg->cfs_bandwidth->runtime来申请；
scheduler_tick() -> task_tick_fair() -> entity_tick() -> update_curr() -> account_cfs_rq_runtime()

```

• scheduler_tick() -> task_tick_fair() -> entity_tick() ->
update_curr() -> account_cfs_rq_runtime()
•
• ↓
•
• static __always_inline
• void account_cfs_rq_runtime(struct cfs_rq *cfs_rq, u64
delta_exec)
• {
•     if (!cfs_bandwidth_used() || !cfs_rq->runtime_enabled)
•         return;
•
•     __account_cfs_rq_runtime(cfs_rq, delta_exec);
• }

```

```

• }
•
• |→
•
• static void __account_cfs_rq_runtime(struct cfs_rq *cfs_rq,
• u64 delta_exec)
• {
•     /* (1) 用cfs_rq已经申请的时间配额(cfs_rq->runtime_remaining)
•     减去已经消耗的时间 */
•     /* dock delta_exec before expiring quota (as it could
•     span periods) */
•     cfs_rq->runtime_remaining -= delta_exec;
•
•     /* (2) cfs_b与cfs_rq的 runtime_expire的比较之后做出决策 */
•     expire_cfs_rq_runtime(cfs_rq);
•
•     /* (3) 如果cfs_rq已经申请的时间配额还没用完, 返回 */
•     if (likely(cfs_rq->runtime_remaining > 0))
•         return;
•
•     /*
•     * if we're unable to extend our runtime we resched so
•     that the active
•     * hierarchy can be throttled
•     */
•     /* (4) 如果cfs_rq申请的时间配额已经用完, 尝试向tg的
•     cfs_b->runtime申请新的时间片
•     如果申请新时间片失败, 说明整个tg已经没有可运行时间了, 把本进程设
•     置为需要重新调度,
•     在中断返回, 发起schedule()时, 发现
•     cfs_rq->runtime_remaining<=0, 会调用throttle_cfs_rq()对cfs_rq进
•     行实质的限制
•     */
•     if (!assign_cfs_rq_runtime(cfs_rq) &&
•         likely(cfs_rq->curr))
•         resched_curr(rq_of(cfs_rq));
• }
•
• ||→
•
• static int assign_cfs_rq_runtime(struct cfs_rq *cfs_rq)
• {
•     struct task_group *tg = cfs_rq->tg;
•     struct cfs_bandwidth *cfs_b = tg_cfs_bandwidth(tg);
•     u64 amount = 0, min_amount, expires;
•
•     /* (4.1) cfs_b的分配时间片的默认值是5ms */
•     /* note: this is a positive sum as runtime_remaining <= 0
•     */
•     min_amount = sched_cfs_bandwidth_slice() -
•     cfs_rq->runtime_remaining;

```

```

• raw_spin_lock(&cfs_b->lock);
• if (cfs_b->quota == RUNTIME_INF)
•     /* (4.2) RUNTIME_INF类型, 时间是分配不完的 */
•     amount = min_amount;
• else {
•     start_cfs_bandwidth(cfs_b);
•
•     /* (4.3) 剩余时间cfs_b->runtime减去分配的时间片, runtime
•     - amount目的是告知系统, 我增加了amount数量的配额, 所以
•     runtime需要减去amount, 表示仅仅运行了runtime-amount时间
•     目的还是按照period做判决throttle */
•     if (cfs_b->runtime > 0) {
•         amount = min(cfs_b->runtime, min_amount);
•         cfs_b->runtime -= amount;
•         cfs_b->idle = 0;
•     }
• }
• expires = cfs_b->runtime_expires;
• raw_spin_unlock(&cfs_b->lock);
•
• /* (4.4) 增加分配的时间片赋值给cfs_rq原先的配额 */
• cfs_rq->runtime_remaining += amount;
• /*
•  * we may have advanced our local expiration to account
• for allowed
•  * spread between our sched_clock and the one on which
• runtime was
•  * issued.
•  */
• if ((s64)(expires - cfs_rq->runtime_expires) > 0)
•     cfs_rq->runtime_expires = expires;
•
• /* (4.5) 判断分配时间是否足够? */
• return cfs_rq->runtime_remaining > 0;
• }

```

3. 在enqueue_task_fair()、put_prev_task_fair()、pick_next_task_fair()这几个时刻, 会 check cfs_rq是否已经达到throttle, 如果达到cfs throttle会把cfs_rq dequeue停止运行;

```

• enqueue_task_fair() -> enqueue_entity() -> check_enqueue_throttle() ->
• throttle_cfs_rq()
• put_prev_task_fair() -> put_prev_entity() -> check_cfs_rq_runtime() ->
• throttle_cfs_rq()
• pick_next_task_fair() -> check_cfs_rq_runtime() -> throttle_cfs_rq()
•
•
• /*
•  * When a group wakes up we want to make sure that its quota is not already
•  * expired/exceeded, otherwise it may be allowed to steal additional ticks
• of
•  * runtime as update_curr() throttling can not trigger until it's on-rq.
•  */
• static void check_enqueue_throttle(struct cfs_rq *cfs_rq)

```

```

{
    if (!cfs_bandwidth_used())
        return;
    /*检测进程组上下节点是否throttle, 并做对应的参数update*/
    /* Synchronize hierarchical throttle counter: */
    if (unlikely(!cfs_rq->throttle_uptodate)) {
        struct rq *rq = rq_of(cfs_rq);
        struct cfs_rq *pcfs_rq;
        struct task_group *tg;

        cfs_rq->throttle_uptodate = 1;

        /* Get closest up-to-date node, because leaves go first: */
        for (tg = cfs_rq->tg->parent; tg; tg = tg->parent) {
            pcfs_rq = tg->cfs_rq[cpu_of(rq)];
            if (pcfs_rq->throttle_uptodate)
                break;
        }
        if (tg) {
            cfs_rq->throttle_count = pcfs_rq->throttle_count;
            cfs_rq->throttled_clock_task = rq_clock_task(rq);
        }
    }

    /* an active group must be handled by the update_curr()->put() path */
    if (!cfs_rq->runtime_enabled || cfs_rq->curr)
        return;
    /*如果已经throttle, 则直接返回*/
    /* ensure the group is not already throttled */
    if (cfs_rq_throttled(cfs_rq))
        return;
    /*update last runtime*/
    /* update runtime allocation */
    account_cfs_rq_runtime(cfs_rq, 0);
    /*配额用完, 进行throttle*/
    if (cfs_rq->runtime_remaining <= 0)
        throttle_cfs_rq(cfs_rq);
}

/* conditionally throttle active cfs_rq's from put_prev_entity() */
static bool check_cfs_rq_runtime(struct cfs_rq *cfs_rq)
{
    if (!cfs_bandwidth_used())
        return false;

    /* (2.1) 如果cfs_rq->runtime_remaining还有运行时间, 直接返回 */
    if (likely(!cfs_rq->runtime_enabled || cfs_rq->runtime_remaining > 0))
        return false;

    /*
     * it's possible for a throttled entity to be forced into a running
     * state (e.g. set_curr_task), in this case we're finished.
     */
    /* (2.2) 如果已经throttle, 直接返回 */

```

```

•   if (cfs_rq_throttled(cfs_rq))
•       return true;
•
•   /* (2.3) 已经throttle, 执行throttle动作 */
•   throttle_cfs_rq(cfs_rq);
•   return true;
• }
•
• static void throttle_cfs_rq(struct cfs_rq *cfs_rq)
• {
•     struct rq *rq = rq_of(cfs_rq);
•     struct cfs_bandwidth *cfs_b = tg_cfs_bandwidth(cfs_rq->tg);
•     struct sched_entity *se;
•     long task_delta, dequeue = 1;
•     bool empty;
•
•     se = cfs_rq->tg->se[cpu_of(rq_of(cfs_rq))];
•
•     /* freeze hierarchy runnable averages while throttled */
•     rcu_read_lock();
•     walk_tg_tree_from(cfs_rq->tg, tg_throttle_down, tg_nop, (void *)rq);
•     rcu_read_unlock();
•
•     task_delta = cfs_rq->h_nr_running;
•     for_each_sched_entity(se) {
•         struct cfs_rq *qcfs_rq = cfs_rq_of(se);
•         /* throttled entity or throttle-on-deactivate */
•         if (!se->on_rq)
•             break;
•
•         /* (3.1) throttle的动作1: 将cfs_rq dequeue停止运行 */
•         if (dequeue)
•             dequeue_entity(qcfs_rq, se, DEQUEUE_SLEEP);
•         qcfs_rq->h_nr_running -= task_delta;
•
•         if (qcfs_rq->load.weight)
•             dequeue = 0;
•     }
•
•     if (!se)
•         sub_nr_running(rq, task_delta);
•
•     /* (3.2) throttle的动作2: 将cfs_rq->throttled置位 */
•     cfs_rq->throttled = 1;
•     cfs_rq->throttled_clock = rq_clock(rq);
•     raw_spin_lock(&cfs_b->lock);
•     empty = list_empty(&cfs_b->throttled_cfs_rq);
•
•     /*
•      * Add to the _head_ of the list, so that an already-started
•      * distribute_cfs_runtime will not see us
•      */
•     list_add_rcu(&cfs_rq->throttled_list, &cfs_b->throttled_cfs_rq);

```

```

•    /*
•    * If we're the first throttled task, make sure the bandwidth
•    * timer is running.
•    */
•    if (empty)
•        start_cfs_bandwidth(cfs_b); /*启动定时器throttle检测*/
•
•    raw_spin_unlock(&cfs_b->lock);
• }

```

4.对每一个tg的cfs_b，系统会启动一个周期性定时器cfs_b->period_timer，运行周期为cfs_b->period。主要作用是period到期后检查是否有cfs_rq被throttle，如果被throttle恢复它，并进行新一轮的监控；

```

• sched_cfs_period_timer() -> do_sched_cfs_period_timer()
•
• ↓
•
• static int do_sched_cfs_period_timer(struct cfs_bandwidth
• *cfs_b, int overrun)
• {
•     u64 runtime, runtime_expires;
•     int throttled;
•
•     /* no need to continue the timer with no bandwidth
• constraint */
•     if (cfs_b->quota == RUNTIME_INF)
•         goto out_deactivate;
•
•     throttled = !list_empty(&cfs_b->throttled_cfs_rq);
•     cfs_b->nr_periods += overrun;
•
•     /*
•      * idle depends on !throttled (for the case of a large
• deficit), and if
•      * we're going inactive then everything else can be
• deferred
•      */
•     if (cfs_b->idle && !throttled)
•         goto out_deactivate;
•
•     /* (1) 新周期的开始，给cfs_b->runtime重新赋值为cfs_b->quota
• 并更新runtime_expires = now + ktime_to_ns(cfs_b->period)
• */
•     __refill_cfs_bandwidth_runtime(cfs_b);
•
•     if (!throttled) {
•         /* mark as potentially idle for the upcoming period
• */
•         cfs_b->idle = 1;
•         return 0;
•     }
• }

```



```

•
•    /* account preceding periods in which throttling occurred
• */
•    cfs_b->nr_throttled += overrun;
•
•    runtime_expires = cfs_b->runtime_expires;
•
•    /*
•     * This check is repeated as we are holding onto the new
bandwidth while
•     * we unthrottle. This can potentially race with an
unthrottled group
•     * trying to acquire new bandwidth from the global pool.
This can result
•     * in us over-using our runtime if it is all used during
this loop, but
•     * only by limited amounts in that extreme case.
•     */
•    /* (2) 解除cfs_b->throttled_cfs_rq中所有被throttle住的cfs_rq
• */
•    while (throttled && cfs_b->runtime > 0) {
•        runtime = cfs_b->runtime;
•        raw_spin_unlock(&cfs_b->lock);
•        /* we can't nest cfs_b->lock while distributing
bandwidth */
•        runtime = distribute_cfs_runtime(cfs_b, runtime,
•                                         runtime_expires);
•        raw_spin_lock(&cfs_b->lock);
•
•        throttled = !list_empty(&cfs_b->throttled_cfs_rq);
•
•        cfs_b->runtime -= min(runtime, cfs_b->runtime);
•    }
•
•    /*
•     * While we are ensured activity in the period following
an
•     * unthrottle, this also covers the case in which the new
bandwidth is
•     * insufficient to cover the existing bandwidth deficit.
(Forcing the
•     * timer to remain active while there are any throttled
entities.)
•     */
•    cfs_b->idle = 0;
•
•    return 0;
•
•out_deactivate:
•    return 1;
• }

```

```

•
• |→
•
• static u64 distribute_cfs_runtime(struct cfs_bandwidth
• *cfs_b,
•     u64 remaining, u64 expires)
• {
•     struct cfs_rq *cfs_rq;
•     u64 runtime;
•     u64 starting_runtime = remaining;
•
•     rcu_read_lock();
•     list_for_each_entry_rcu(cfs_rq, &cfs_b->throttled_cfs_rq,
•         throttled_list) {
•         struct rq *rq = rq_of(cfs_rq);
•
•         raw_spin_lock(&rq->lock);
•         if (!cfs_rq_throttled(cfs_rq))
•             goto next;
•
•         runtime = -cfs_rq->runtime_remaining + 1;
•         if (runtime > remaining)
•             runtime = remaining;
•         remaining -= runtime;
•
•         cfs_rq->runtime_remaining += runtime;
•         cfs_rq->runtime_expires = expires;
•
•         /* (2.1) 解除throttle */
•         /* we check whether we're throttled above */
•         if (cfs_rq->runtime_remaining > 0)
•             unthrottle_cfs_rq(cfs_rq);
•
• next:
•         raw_spin_unlock(&rq->lock);
•
•         if (!remaining)
•             break;
•     }
•     rcu_read_unlock();
•
•     return starting_runtime - remaining;
• }
•
• ||→
•
• void unthrottle_cfs_rq(struct cfs_rq *cfs_rq)
• {
•     struct rq *rq = rq_of(cfs_rq);
•     struct cfs_bandwidth *cfs_b =
•         tg_cfs_bandwidth(cfs_rq->tg);

```

```

• struct sched_entity *se;
• int enqueue = 1;
• long task_delta;
•
• se = cfs_rq->tg->se[cpu_of(rq)];
•
• cfs_rq->throttled = 0;
•
• update_rq_clock(rq);
•
• raw_spin_lock(&cfs_b->lock);
• cfs_b->throttled_time += rq_clock(rq) -
cfs_rq->throttled_clock;
• list_del_rcu(&cfs_rq->throttled_list);
• raw_spin_unlock(&cfs_b->lock);
•
• /* update hierarchical throttle state */
• walk_tg_tree_from(cfs_rq->tg, tg_nop, tg_unthrottle_up,
(void *)rq);
•
• if (!cfs_rq->load.weight)
•     return;
•
• task_delta = cfs_rq->h_nr_running;
• for_each_sched_entity(se) {
•     if (se->on_rq)
•         enqueue = 0;
•
•     cfs_rq = cfs_rq_of(se);
•     /* (2.1.1) 重新enqueue运行 */
•     if (enqueue)
•         enqueue_entity(cfs_rq, se, ENQUEUE_WAKEUP);
•     cfs_rq->h_nr_running += task_delta;
•
•     if (cfs_rq_throttled(cfs_rq))
•         break;
• }
•
• if (!se)
•     add_nr_running(rq, task_delta);
•
• /* determine whether we need to wake up potentially idle
cpu */
• if (rq->curr == rq->idle && rq->cfs.nr_running)
•     resched_curr(rq);
• }

```

明白其思路就可以.但是我看了好几个手机平台都没有定义CONFIG_CFS_BANDWIDTH,似乎都没有使用.目前cpu速度越来越快,处理能力一般都没什么问题,不需要throttle.

