

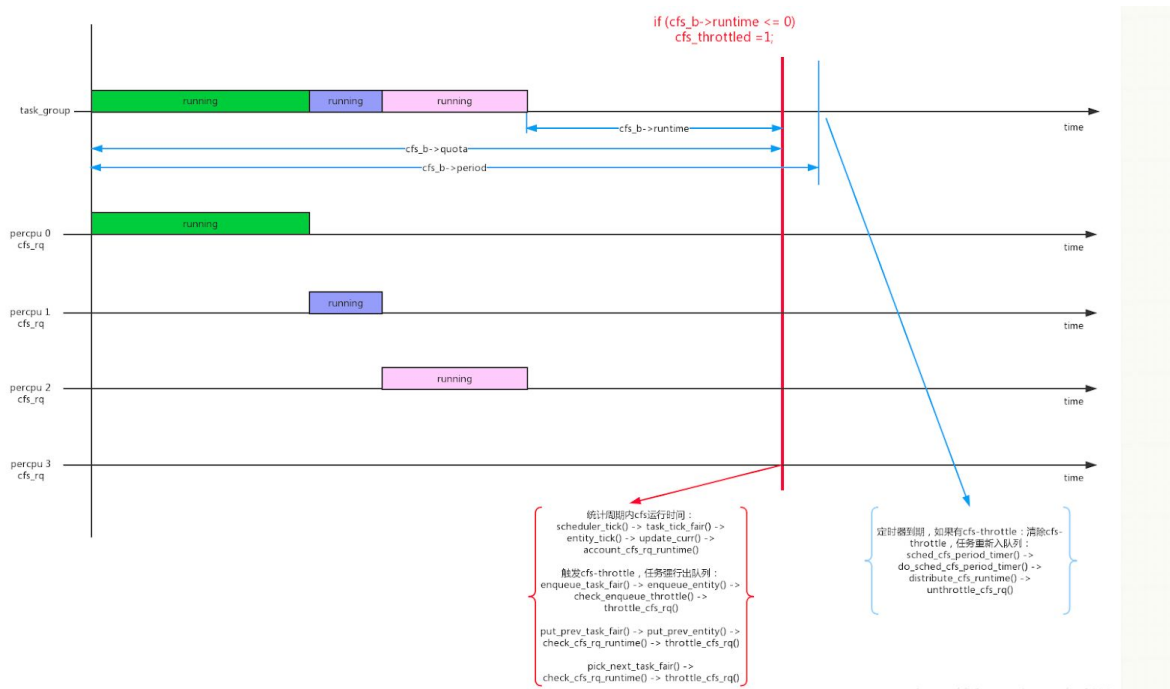
# cfs bandwidth

1 cfs bandwidth 概述	1
2. cfs bandwidth 额度分配	4
3. 何时进行bandwidth throttle	6
4. 已经throttle的cfs_rq, 如何解除	8

## 1 cfs bandwidth 概述

cfs bandwidth是针对task\_group的配置, 一个task\_group的bandwidth使用一个struct cfs\_bandwidth \*cfs\_b数据结构来控制。

```
• struct cfs_bandwidth {
• #ifdef CONFIG_CFS_BANDWIDTH
•     raw_spinlock_t lock;
•     /*cfs bandwidth的监控周期, 默认值是default_cfs_period() 0.1s
•     */
•     ktime_t period;
•     /* quota:cfs task_group 在一个监控周期内的运行时间配额, 默认值是RUNTIME_INF,
•     无限大
•
•     runtime:cfs task_group 在一个监控周期内剩余可运行的时间*/
•     u64 quota, runtime;
•     s64 hierarchical_quota;
•     u64 runtime_expires;
•
•     int idle, period_active;
•     /*period_timer周期性的throttle动作,slack_timer是idle时候的timer*/
•     struct hrtimer period_timer, slack_timer;
•     struct list_head throttled_cfs_rq;
•
•     /* statistics */
•     int nr_periods, nr_throttled;
•     u64 throttled_time;
• #endif
• };
```



我们首先通过运行图来简单的了解其工作原理:

- 系统首先会预算一个运行时间配额和剩余运行时间,两者默认是相等的
- 当某个task\_group里的task开始运行一段时间之后,比如为delta,则剩余运行时间变成了初始的剩余运行时间-delta,更新新的剩余运行时间
- 如果在一个周期里面,剩余运行时间用光了,可以尝试那补偿5ms的时间,即总的运行时间减少了5ms,而剩余运行时间增加了5ms.
- 随着时间的流逝,剩余运行时间逐渐减少到0甚至为负值,如果在检测过程中,检测到了剩余运行时间已经使用完毕,那么系统就会额外的补偿给剩余运行时间数值为 5-runtime\_remaining(unit:ms).
- 在每次pick task的时候都会检测是否可以throttle,如果可以,则强制将enqueue的task dequeue,并有一个period timer(100ms)定时检测是否有rq throttle了,如果有则cfs调度算法重新对task进行调度操作.

下面是它的初始化:

```

/*执行slack_timer的回调函数*/
static enum hrtimer_restart sched_cfs_slack_timer(struct hrtimer *timer)
{
    struct cfs_bandwidth *cfs_b =
        container_of(timer, struct cfs_bandwidth, slack_timer);

    do_sched_cfs_slack_timer(cfs_b);

    return HRTIMER_NORESTART;
}

/*running period timer*/
static enum hrtimer_restart sched_cfs_period_timer(struct hrtimer *timer)
{
    struct cfs_bandwidth *cfs_b =
        container_of(timer, struct cfs_bandwidth, period_timer);

    int overrun;

    int idle = 0;

```

```

•
• raw_spin_lock(&cfs_b->lock);
• for (;;) {
•     overrun = hrtimer_forward_now(timer, cfs_b->period);
•     if (!overrun)
•         break;
•
•     idle = do_sched_cfs_period_timer(cfs_b, overrun);
• }
• if (idle)
•     cfs_b->period_active = 0;
• raw_spin_unlock(&cfs_b->lock);
•
• return idle ? HRTIMER_NORESTART : HRTIMER_RESTART;
• }
• /*
•  * default period for cfs group bandwidth.
•  * default: 0.1s, units: nanoseconds
•  */
• static inline u64 default_cfs_period(void)
• {
•     return 100000000ULL;
• }
• /*cfs bandwidth的初始化*/
• void init_cfs_bandwidth(struct cfs_bandwidth *cfs_b)
• {
•     raw_spin_lock_init(&cfs_b->lock);
•     cfs_b->runtime = 0;
•     cfs_b->quota = RUNTIME_INF;
•     cfs_b->period = ns_to_ktime(default_cfs_period());
•
•     INIT_LIST_HEAD(&cfs_b->throttled_cfs_rq);
•     /*周期性处理cfs bandwidth上的task_group*/
•     hrtimer_init(&cfs_b->period_timer, CLOCK_MONOTONIC,
HRTIMER_MODE_ABS_PINNED);
•     cfs_b->period_timer.function = sched_cfs_period_timer;
•     hrtimer_init(&cfs_b->slack_timer, CLOCK_MONOTONIC, HRTIMER_MODE_REL);
•     cfs_b->slack_timer.function = sched_cfs_slack_timer;
• }
•
• static void init_cfs_rq_runtime(struct cfs_rq *cfs_rq)
• {
•     cfs_rq->runtime_enabled = 0;
•     INIT_LIST_HEAD(&cfs_rq->throttled_list);
• }
•
• void start_cfs_bandwidth(struct cfs_bandwidth *cfs_b)
• {
•     lockdep_assert_held(&cfs_b->lock);
•
•     if (!cfs_b->period_active) {
•         cfs_b->period_active = 1;
•         hrtimer_forward_now(&cfs_b->period_timer, cfs_b->period);
•     }
• }

```

```

•     hrtimer_start_expires(&cfs_b->period_timer,
•     HRTIMER_MODE_ABS_PINNED);
•     }
•     }

```

## 2. cfs bandwidth 额度分配

因为一个task\_group是在percpu上都创建了一个cfs\_rq，所以cfs\_b->quota的值是这些percpu cfs\_rq中的进程共享的，每个percpu cfs\_rq在运行时需要向tg->cfs\_bandwidth->runtime来申请；

scheduler\_tick() -> task\_tick\_fair() -> entity\_tick() -> update\_curr() -> account\_cfs\_rq\_runtime()

```

• scheduler_tick() -> task_tick_fair() -> entity_tick() ->
• update_curr() -> account_cfs_rq_runtime()
•
• ↓
•
• static __always_inline
• void account_cfs_rq_runtime(struct cfs_rq *cfs_rq, u64
• delta_exec)
• {
•     if (!cfs_bandwidth_used() || !cfs_rq->runtime_enabled)
•         return;
•
•     __account_cfs_rq_runtime(cfs_rq, delta_exec);
• }
•
• |→
•
• static void __account_cfs_rq_runtime(struct cfs_rq *cfs_rq,
• u64 delta_exec)
• {
•     /* (1) 用cfs_rq已经申请的时间配额(cfs_rq->runtime_remaining)
• 减去已经消耗的时间 */
•     /* dock delta_exec before expiring quota (as it could
• span periods) */
•     cfs_rq->runtime_remaining -= delta_exec;
•
•     /* (2) cfs_b与cfs_rq的 runtime_expire的比较之后做出决策 */
•     expire_cfs_rq_runtime(cfs_rq);
•
•     /* (3) 如果cfs_rq已经申请的时间配额还没用完, 返回 */
•     if (likely(cfs_rq->runtime_remaining > 0))
•         return;
•
•     /*
•      * if we're unable to extend our runtime we resched so
•      that the active

```

```

•      * hierarchy can be throttled
•      */
•      /* (4) 如果cfs_rq申请的时间配额已经用完, 尝试向tg的
cfs_b->runtime申请新的时间片
•      如果申请新时间片失败, 说明整个tg已经没有可运行时间了, 把本进程设
置为需要重新调度,
•      在中断返回, 发起schedule()时, 发现
cfs_rq->runtime_remaining<=0, 会调用throttle_cfs_rq()对cfs_rq进行
实质的限制
•      */
•      if (!assign_cfs_rq_runtime(cfs_rq) &&
likely(cfs_rq->curr))
•          resched_curr(rq_of(cfs_rq));
•      }
•
•      ||→
•
•      static int assign_cfs_rq_runtime(struct cfs_rq *cfs_rq)
•      {
•          struct task_group *tg = cfs_rq->tg;
•          struct cfs_bandwidth *cfs_b = tg_cfs_bandwidth(tg);
•          u64 amount = 0, min_amount, expires;
•
•          /* (4.1) cfs_b的分配时间片的默认值是5ms */
•          /* note: this is a positive sum as runtime_remaining <= 0
*/
•          min_amount = sched_cfs_bandwidth_slice() -
cfs_rq->runtime_remaining;
•
•          raw_spin_lock(&cfs_b->lock);
•          if (cfs_b->quota == RUNTIME_INF)
•              /* (4.2) RUNTIME_INF类型, 时间是分配不完的 */
•              amount = min_amount;
•          else {
•              start_cfs_bandwidth(cfs_b);
•
•              /* (4.3) 剩余时间cfs_b->runtime减去分配的时间片, runtime
- amount目的是告知系统, 我增加了amount数量的配额, 所以
runtime需要减去amount, 表示仅仅运行了runtime-amount时间
目的还是按照period做判决throttle */
•              if (cfs_b->runtime > 0) {
•                  amount = min(cfs_b->runtime, min_amount);
•                  cfs_b->runtime -= amount;
•                  cfs_b->idle = 0;
•              }
•          }
•          expires = cfs_b->runtime_expires;
•          raw_spin_unlock(&cfs_b->lock);
•
•          /* (4.4) 增加分配的时间片赋值给cfs_rq原先的配额 */
•          cfs_rq->runtime_remaining += amount;

```

```

•      /*
•      * we may have advanced our local expiration to account
•      for allowed
•      * spread between our sched_clock and the one on which
•      runtime was
•      * issued.
•      */
•      if ((s64)(expires - cfs_rq->runtime_expires) > 0)
•          cfs_rq->runtime_expires = expires;
•
•      /* (4.5) 判断分配时间是否足够? */
•      return cfs_rq->runtime_remaining > 0;
•  }

```

### 3. 何时进行bandwidth throttle

在enqueue\_task\_fair()、put\_prev\_task\_fair()、pick\_next\_task\_fair()这几个时刻，会 check cfs\_rq是否已经达到throttle，如果达到cfs throttle会把cfs\_rq dequeue停止运行；

```

•  enqueue_task_fair() -> enqueue_entity() -> check_enqueue_throttle() ->
•  throttle_cfs_rq()
•  put_prev_task_fair() -> put_prev_entity() -> check_cfs_rq_runtime() ->
•  throttle_cfs_rq()
•  pick_next_task_fair() -> check_cfs_rq_runtime() -> throttle_cfs_rq()
•
•
•
•  /*
•  * When a group wakes up we want to make sure that its quota is not already
•  * expired/exceeded, otherwise it may be allowed to steal additional ticks
•  of
•  * runtime as update_curr() throttling can not trigger until it's on-rq.
•  */
•  static void check_enqueue_throttle(struct cfs_rq *cfs_rq)
•  {
•      if (!cfs_bandwidth_used())
•          return;
•      /*检测进程组上下节点是否throttle, 并做对应的参数update*/
•      /* Synchronize hierarchical throttle counter: */
•      if (unlikely(!cfs_rq->throttle_uptodate)) {
•          struct rq *rq = rq_of(cfs_rq);
•          struct cfs_rq *pcfs_rq;
•          struct task_group *tg;
•
•          cfs_rq->throttle_uptodate = 1;
•
•          /* Get closest up-to-date node, because leaves go first: */
•          for (tg = cfs_rq->tg->parent; tg; tg = tg->parent) {
•              pcfs_rq = tg->cfs_rq[cpu_of(rq)];
•              if (pcfs_rq->throttle_uptodate)
•                  break;
•          }
•          if (tg) {

```

```

    cfs_rq->throttle_count = pcfs_rq->throttle_count;
    cfs_rq->throttled_clock_task = rq_clock_task(rq);
}

/* an active group must be handled by the update_curr()->put() path */
if (!cfs_rq->runtime_enabled || cfs_rq->curr)
    return;
/* 如果已经throttle,则直接返回 */
/* ensure the group is not already throttled */
if (cfs_rq_throttled(cfs_rq))
    return;
/* update last runtime */
/* update runtime allocation */
account_cfs_rq_runtime(cfs_rq, 0);
/* 配额用完,进行throttle */
if (cfs_rq->runtime_remaining <= 0)
    throttle_cfs_rq(cfs_rq);
}

/* conditionally throttle active cfs_rq's from put_prev_entity() */
static bool check_cfs_rq_runtime(struct cfs_rq *cfs_rq)
{
    if (!cfs_bandwidth_used())
        return false;

    /* (2.1) 如果cfs_rq->runtime_remaining还有运行时间, 直接返回 */
    if (likely(!cfs_rq->runtime_enabled || cfs_rq->runtime_remaining > 0))
        return false;

    /*
     * it's possible for a throttled entity to be forced into a running
     * state (e.g. set_curr_task), in this case we're finished.
     */
    /* (2.2) 如果已经throttle, 直接返回 */
    if (cfs_rq_throttled(cfs_rq))
        return true;

    /* (2.3) 已经throttle, 执行throttle动作 */
    throttle_cfs_rq(cfs_rq);
    return true;
}

static void throttle_cfs_rq(struct cfs_rq *cfs_rq)
{
    struct rq *rq = rq_of(cfs_rq);
    struct cfs_bandwidth *cfs_b = tg_cfs_bandwidth(cfs_rq->tg);
    struct sched_entity *se;
    long task_delta, dequeue = 1;
    bool empty;

    se = cfs_rq->tg->se[cpu_of(rq_of(cfs_rq))];

    /* freeze hierarchy runnable averages while throttled */

```

```

• rcu_read_lock();
• walk_tg_tree_from(cfs_rq->tg, tg_throttle_down, tg_nop, (void *)rq);
• rcu_read_unlock();
•
• task_delta = cfs_rq->h_nr_running;
• for_each_sched_entity(se) {
•     struct cfs_rq *qcfs_rq = cfs_rq_of(se);
•     /* throttled entity or throttle-on-deactivate */
•     if (!se->on_rq)
•         break;
•
•     /* (3.1) throttle的动作1: 将cfs_rq dequeue停止运行 */
•     if (dequeue)
•         dequeue_entity(qcfs_rq, se, DEQUEUE_SLEEP);
•     qcfs_rq->h_nr_running -= task_delta;
•
•     if (qcfs_rq->load.weight)
•         dequeue = 0;
• }
•
• if (!se)
•     sub_nr_running(rq, task_delta);
•
• /* (3.2) throttle的动作2: 将cfs_rq->throttled置位 */
• cfs_rq->throttled = 1;
• cfs_rq->throttled_clock = rq_clock(rq);
• raw_spin_lock(&cfs_b->lock);
• empty = list_empty(&cfs_b->throttled_cfs_rq);
•
• /*
•  * Add to the _head_ of the list, so that an already-started
•  * distribute_cfs_runtime will not see us
•  */
• list_add_rcu(&cfs_rq->throttled_list, &cfs_b->throttled_cfs_rq);
•
• /*
•  * If we're the first throttled task, make sure the bandwidth
•  * timer is running.
•  */
• if (empty)
•     start_cfs_bandwidth(cfs_b); /*启动定时器throttle检测*/
•
• raw_spin_unlock(&cfs_b->lock);
• }

```

## 4. 已经throttle的cfs\_rq，如何解除

对每一个tg的cfs\_b，系统会启动一个周期性定时器cfs\_b->period\_timer，运行周期为cfs\_b->period。主要作用是period到期后检查是否有cfs\_rq被throttle，如果被throttle恢复它，并进行新一轮的监控；

- sched\_cfs\_period\_timer() -> do\_sched\_cfs\_period\_timer()



```

•
• ↓
•
• static int do_sched_cfs_period_timer(struct cfs_bandwidth
• *cfs_b, int overrun)
• {
•     u64 runtime, runtime_expires;
•     int throttled;
•
•     /* no need to continue the timer with no bandwidth
• constraint */
•     if (cfs_b->quota == RUNTIME_INF)
•         goto out_deactivate;
•
•     throttled = !list_empty(&cfs_b->throttled_cfs_rq);
•     cfs_b->nr_periods += overrun;
•
•     /*
•      * idle depends on !throttled (for the case of a large
• deficit), and if
•      * we're going inactive then everything else can be
• deferred
•      */
•     if (cfs_b->idle && !throttled)
•         goto out_deactivate;
•
•     /* (1) 新周期的开始, 给cfs_b->runtime重新赋值为cfs_b->quota
• 并更新runtime_expires = now + ktime_to_ns(cfs_b->period)
• */
•     __refill_cfs_bandwidth_runtime(cfs_b);
•
•     if (!throttled) {
•         /* mark as potentially idle for the upcoming period
• */
•         cfs_b->idle = 1;
•         return 0;
•     }
•
•     /* account preceding periods in which throttling occurred
• */
•     cfs_b->nr_throttled += overrun;
•
•     runtime_expires = cfs_b->runtime_expires;
•
•     /*
•      * This check is repeated as we are holding onto the new
• bandwidth while
•      * we unthrottle. This can potentially race with an
• unthrottled group
•      * trying to acquire new bandwidth from the global pool.
• This can result

```

```

•      * in us over-using our runtime if it is all used during
this loop, but
•      * only by limited amounts in that extreme case.
•      */
•      /* (2) 解除cfs_b->throttled_cfs_rq中所有被throttle住的cfs_rq
*/
•      while (throttled && cfs_b->runtime > 0) {
•          runtime = cfs_b->runtime;
•          raw_spin_unlock(&cfs_b->lock);
•          /* we can't nest cfs_b->lock while distributing
bandwidth */
•          runtime = distribute_cfs_runtime(cfs_b, runtime,
•                                          runtime_expires);
•          raw_spin_lock(&cfs_b->lock);
•
•          throttled = !list_empty(&cfs_b->throttled_cfs_rq);
•
•          cfs_b->runtime -= min(runtime, cfs_b->runtime);
•      }
•
•      /*
•      * While we are ensured activity in the period following
an
•      * unthrottle, this also covers the case in which the new
bandwidth is
•      * insufficient to cover the existing bandwidth deficit.
(Forcing the
•      * timer to remain active while there are any throttled
entities.)
•      */
•      cfs_b->idle = 0;
•
•      return 0;
•
• out_deactivate:
•      return 1;
•  }
•
•  |→
•
•  static u64 distribute_cfs_runtime(struct cfs_bandwidth
•  *cfs_b,
•      u64 remaining, u64 expires)
•  {
•      struct cfs_rq *cfs_rq;
•      u64 runtime;
•      u64 starting_runtime = remaining;
•
•      rcu_read_lock();
•      list_for_each_entry_rcu(cfs_rq, &cfs_b->throttled_cfs_rq,
•                              throttled_list) {

```

```

•      struct rq *rq = rq_of(cfs_rq);
•
•      raw_spin_lock(&rq->lock);
•      if (!cfs_rq_throttled(cfs_rq))
•          goto next;
•
•      runtime = -cfs_rq->runtime_remaining + 1;
•      if (runtime > remaining)
•          runtime = remaining;
•      remaining -= runtime;
•
•      cfs_rq->runtime_remaining += runtime;
•      cfs_rq->runtime_expires = expires;
•
•      /* (2.1) 解除throttle */
•      /* we check whether we're throttled above */
•      if (cfs_rq->runtime_remaining > 0)
•          unthrottle_cfs_rq(cfs_rq);
•
• next:
•      raw_spin_unlock(&rq->lock);
•
•      if (!remaining)
•          break;
•  }
•  rcu_read_unlock();
•
•  return starting_runtime - remaining;
• }
•
• ||→
•
• void unthrottle_cfs_rq(struct cfs_rq *cfs_rq)
• {
•     struct rq *rq = rq_of(cfs_rq);
•     struct cfs_bandwidth *cfs_b =
tg_cfs_bandwidth(cfs_rq->tg);
•     struct sched_entity *se;
•     int enqueue = 1;
•     long task_delta;
•
•     se = cfs_rq->tg->se[cpu_of(rq)];
•
•     cfs_rq->throttled = 0;
•
•     update_rq_clock(rq);
•
•     raw_spin_lock(&cfs_b->lock);
•     cfs_b->throttled_time += rq_clock(rq) -
cfs_rq->throttled_clock;
•     list_del_rcu(&cfs_rq->throttled_list);

```

```

• raw_spin_unlock(&cfs_b->lock);
•
• /* update hierarchical throttle state */
• walk_tg_tree_from(cfs_rq->tg, tg_nop, tg_unthrottle_up,
• (void *)rq);
•
• if (!cfs_rq->load.weight)
•     return;
•
• task_delta = cfs_rq->h_nr_running;
• for_each_sched_entity(se) {
•     if (se->on_rq)
•         enqueue = 0;
•
•     cfs_rq = cfs_rq_of(se);
•     /* (2.1.1) 重新enqueue运行 */
•     if (enqueue)
•         enqueue_entity(cfs_rq, se, ENQUEUE_WAKEUP);
•     cfs_rq->h_nr_running += task_delta;
•
•     if (cfs_rq_throttled(cfs_rq))
•         break;
• }
•
• if (!se)
•     add_nr_running(rq, task_delta);
•
• /* determine whether we need to wake up potentially idle
•  cpu */
• if (rq->curr == rq->idle && rq->cfs.nr_running)
•     resched_curr(rq);
• }

```

明白其思路就可以.但是我看了好几个手机平台都没有定义CONFIG\_CFS\_BANDWIDTH,似乎都没有使用.目前cpu速度越来越快,处理能力一般都没什么问题,不需要throttle.



