

由于在task在创建或者task从idle会wakeup的流程中涉及到sched domain和sched group的知识,所以现在提前看下这部分.

根据实际的物理属性,CPU domain分成下面三种:

CPU 分类	Linux内核分类	描述
SMT(多线程)	CONFIG_SHCED_SMT	一个物理核心可以有两个以及之上的执行线程,一个物理核心里面的线程共享相同的CPU资源和L1 cache,task的迁移不会影响cache利用率
MC(多核)	CONFIG_SHCED_MC	每个物理核心有独立的L1 cache,多个物理核心可以组成一个cluster,cluster共享L2 cache
DIE(SoC)	简称为DIE	SoC级别

## 一 cpu topology的获取

arm64架构的cpu拓扑结构存储在cpu\_topology[NR\_CPUS]变量当中 :

```
• struct cpu_topology {
•     int thread_id;
•     int core_id;
•     int cluster_id;
•     cpumask_t thread_sibling;
•     cpumask_t core_sibling;
• };
•
• extern struct cpu_topology cpu_topology[NR_CPUS];
• #define topology_physical_package_id(cpu) (cpu_topology[cpu].cluster_id)
• #define topology_core_id(cpu) (cpu_topology[cpu].core_id)
• #define topology_core_cpumask(cpu) (&cpu_topology[cpu].core_sibling)
• #define topology_sibling_cpumask(cpu) (&cpu_topology[cpu].thread_sibling)
```

cpu\_topology[]数组获取的路径如下(都是从dts里面获取的):

```
• kernel_init() -> kernel_init_freeable() -> smp_prepare_cpus() ->
• init_cpu_topology() -> parse_dt_topology()
• --->
• static int __init parse_dt_topology(void)
• {
•     struct device_node *cn, *map;
•     int ret = 0;
•     int cpu;
•     /*从dts里面获取cpu node的跟节点*/
•     cn = of_find_node_by_path("/cpus");
•     if (!cn) {
•         pr_err("No CPU information found in DT\n");
•         return 0;
•     }
•
•     /*
•      * When topology is provided cpu-map is essentially a root
•      * cluster with restricted subnodes.
```

```

    */
    map = of_get_child_by_name(cn, "cpu-map");
    if (!map)
        goto out;

    ret = parse_cluster(map, 0);
    if (ret != 0)
        goto out_map;

    /*
     * Check that all cores are in the topology; the SMP code will
     * only mark cores described in the DT as possible.
     */
    for_each_possible_cpu(cpu)
        if (cpu_topology[cpu].cluster_id == -1)
            ret = -EINVAL;

out_map:
    of_node_put(map);
out:
    of_node_put(cn);
    return ret;
}
/*看下解析cluster怎么做的*/
static int __init parse_cluster(struct device_node *cluster, int depth)
{
    char name[10];
    bool leaf = true;
    bool has_cores = false;
    struct device_node *c;
    static int cluster_id __initdata;
    int core_id = 0;
    int i, ret;

    /*
     * First check for child clusters; we currently ignore any
     * information about the nesting of clusters and present the
     * scheduler with a flat list of them.
     */
    i = 0;
    do { /*多个cluster,迭代遍历*/
        snprintf(name, sizeof(name), "cluster%d", i);
        c = of_get_child_by_name(cluster, name);
        if (c) {
            leaf = false;
            ret = parse_cluster(c, depth + 1);
            of_node_put(c);
            if (ret != 0)
                return ret;
        }
        i++;
    } while (c);

    /* Now check for cores */
    i = 0;
    do { /*遍历cluster下面core的节点*/
        snprintf(name, sizeof(name), "core%d", i);
        c = of_get_child_by_name(cluster, name);
        if (c) {
            has_cores = true;

```

```

        if (depth == 0) {
            pr_err("%s: cpu-map children should be clusters\n",
                   c->full_name);
            of_node_put(c);
            return -EINVAL;
        }
        /*如果是叶子cluster节点，继续遍历core中的cpu节点*/
        if (leaf) {
            ret = parse_core(c, cluster_id, core_id++);
        } else {
            pr_err("%s: Non-leaf cluster with core %s\n",
                   cluster->full_name, name);
            ret = -EINVAL;
        }

        of_node_put(c);
        if (ret != 0)
            return ret;
    }
    i++;
} while (c);

if (leaf && !has_cores)
    pr_warn("%s: empty cluster\n", cluster->full_name);

if (leaf)
    cluster_id++;

return 0;
}

static int __init parse_core(struct device_node *core, int cluster_id,
                             int core_id)
{
    char name[10];
    bool leaf = true;
    int i = 0;
    int cpu;
    struct device_node *t;

    do { /*如果存在thread层级，解析thread和cpu层级*/
        snprintf(name, sizeof(name), "thread%d", i);
        t = of_get_child_by_name(core, name);
        if (t) {
            leaf = false;
            cpu = get_cpu_for_node(t);
            if (cpu >= 0) {
                cpu_topology[cpu].cluster_id = cluster_id;
                cpu_topology[cpu].core_id = core_id;
                cpu_topology[cpu].thread_id = i;
            } else {
                pr_err("%s: Can't get CPU for thread\n",
                       t->full_name);
                of_node_put(t);
                return -EINVAL;
            }
            of_node_put(t);
        }
        i++;
    } while (t);
    /*否则直接解析cpu层级*/
    cpu = get_cpu_for_node(core);
    if (cpu >= 0) {
        if (!leaf) {
            pr_err("%s: Core has both threads and CPU\n",
                   core->full_name);
            return -EINVAL;
        }
    }
}

```

```

    }
    /*得到了cpu的cluster_id/core_id*/
    cpu_topology[cpu].cluster_id = cluster_id;
    cpu_topology[cpu].core_id = core_id;
} else if (leaf) {
    pr_err("%s: Can't get CPU for leaf core\n", core->full_name);
    return -EINVAL;
}

return 0;
}

static int __init get_cpu_for_node(struct device_node *node)
{
    struct device_node *cpu_node;
    int cpu;
    /*解析cpu层级*/
    cpu_node = of_parse_phandle(node, "cpu", 0);
    if (!cpu_node)
        return -1;

    for_each_possible_cpu(cpu) {
        if (of_get_cpu_node(cpu, NULL) == cpu_node) {
            of_node_put(cpu_node);
            return cpu;
        }
    }

    pr_crit("Unable to find CPU node for %s\n", cpu_node->full_name);

    of_node_put(cpu_node);
    return -1;
}

```

cpu同一层次的关系cpu\_topology[cpu].core\_sibling/thread\_sibling会在update\_siblings\_masks()中更新

```

kernel_init() -> kernel_init_freeable() -> smp_prepare_cpus()
-> store_cpu_topology() -> update_siblings_masks()
--->
static void update_siblings_masks(unsigned int cpuid)
{
    struct cpu_topology *cpu_topo, *cpuid_topo =
    &cpu_topology[cpuid];
    int cpu;

    /* update core and thread sibling masks */
    for_each_possible_cpu(cpu) {
        cpu_topo = &cpu_topology[cpu];

        if (cpuid_topo->cluster_id != cpu_topo->cluster_id)
            continue;

        cpumask_set_cpu(cpuid, &cpu_topo->core_sibling);
        if (cpu != cpuid)
            cpumask_set_cpu(cpu, &cpuid_topo->core_sibling);

        if (cpuid_topo->core_id != cpu_topo->core_id)
            continue;
    }
}

```

```

•         cpumask_set_cpu(cpuid, &cpu_topo->thread_sibling);
•         if (cpu != cpuid)
•             cpumask_set_cpu(cpu,
• &cpuid_topo->thread_sibling);
•     }
• }

```

以我现在的手机平台为例,cpu相关的dts如下:

```

•     cpus {
•         #address-cells = <2>;
•         #size-cells = <0>;
•
•         cpu-map {
•             cluster0 {
•                 core0 {
•                     cpu = <&CPU0>;
•                 };
•                 core1 {
•                     cpu = <&CPU1>;
•                 };
•                 core2 {
•                     cpu = <&CPU2>;
•                 };
•                 core3 {
•                     cpu = <&CPU3>;
•                 };
•             };
•             cluster1 {
•                 core0 {
•                     cpu = <&CPU4>;
•                 };
•                 core1 {
•                     cpu = <&CPU5>;
•                 };
•                 core2 {
•                     cpu = <&CPU6>;
•                 };
•                 core3 {
•                     cpu = <&CPU7>;
•                 };
•             };
•         };
•
•         CPU0: cpu@0 {
•             device_type = "cpu";
•             compatible = "arm,cortex-a55", "arm,armv8";
•             reg = <0x0 0x0>;
•             enable-method = "psci";
•             cpu-supply = <&fan53555_dcdc>;
•             cpufreq-data-v1 = <&cpufreq_clus0>;
•             cpu-idle-states = <&CORE_PD>;
•             sched-energy-costs = <&CPU_COST_0 &CLUSTER_COST_0>;
•         };
•         CPU1: cpu@100 {
•             device_type = "cpu";

```

```

•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x100>;
•         enable-method = "psci";
•         cpu-supply = <&fan53555_dcdc>;
•         cpufreq-data-v1 = <&cpufreq_clus0>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_0 &CLUSTER_COST_0>;
•     };
• CPU2: cpu@200 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x200>;
•         enable-method = "psci";
•         cpu-supply = <&fan53555_dcdc>;
•         cpufreq-data-v1 = <&cpufreq_clus0>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_0 &CLUSTER_COST_0>;
•     };
• CPU3: cpu@300 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x300>;
•         enable-method = "psci";
•         cpu-supply = <&fan53555_dcdc>;
•         cpufreq-data-v1 = <&cpufreq_clus0>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_0 &CLUSTER_COST_0>;
•     };
• CPU4: cpu@400 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x400>;
•         enable-method = "psci";
•         cpu-supply = <&vddcpu>;
•         cpufreq-data-v1 = <&cpufreq_clus1>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_1 &CLUSTER_COST_1>;
•     };
• CPU5: cpu@500 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x500>;
•         enable-method = "psci";
•         cpu-supply = <&vddcpu>;
•         cpufreq-data-v1 = <&cpufreq_clus1>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_1 &CLUSTER_COST_1>;
•     };
• CPU6: cpu@600 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55", "arm,armv8";
•         reg = <0x0 0x600>;
•         enable-method = "psci";
•         cpu-supply = <&vddcpu>;

```

```

•         cpufreq-data-v1 = <&cpufreq_clus1>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_1 &CLUSTER_COST_1>;
•     };
•     CPU7: cpu@700 {
•         device_type = "cpu";
•         compatible = "arm,cortex-a55","arm,armv8";
•         reg = <0x0 0x700>;
•         enable-method = "psci";
•         cpu-supply = <&vddcpu>;
•         cpufreq-data-v1 = <&cpufreq_clus1>;
•         cpu-idle-states = <&CORE_PD>;
•         sched-energy-costs = <&CPU_COST_1 &CLUSTER_COST_1>;
•     };
• };

```

经过parse\_dt\_topology()、update\_siblings\_masks()解析后得到cpu\_topology[]的值为：

```

•  cpu 0 cluster_id = 0, core_id = 0, core_sibling = 0xf
•  cpu 1 cluster_id = 0, core_id = 1, core_sibling = 0xf
•  cpu 2 cluster_id = 0, core_id = 2, core_sibling = 0xf
•  cpu 3 cluster_id = 0, core_id = 3, core_sibling = 0xf
•  cpu 4 cluster_id = 1, core_id = 0, core_sibling = 0xf0
•  cpu 5 cluster_id = 1, core_id = 1, core_sibling = 0xf0
•  cpu 6 cluster_id = 1, core_id = 2, core_sibling = 0xf0
•  cpu 7 cluster_id = 1, core_id = 3, core_sibling = 0xf0
•  //adb 查看cpu topology信息
•  meizul000:/sys/devices/system/cpu/cpu0/topology # ls -l
•  total 0
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 core_id //0
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 core_siblings //0f
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 core_siblings_list//0-3
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 physical_package_id //0
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 thread_siblings //01
•  -r--r--r-- 1 root root 4096 2018-09-05 16:10 thread_siblings_list//0
•

```

## 二 调度域,调度组以及调度组capacity的初始化

内核中有预先定义的调度域的结构体struct sched\_domain\_topology\_level:

```

•  typedef const struct cpumask *(*sched_domain_mask_f)(int cpu);
•  typedef int (*sched_domain_flags_f)(void);
•  typedef
•  const struct sched_group_energy * const(*sched_domain_energy_f)(int cpu);
•
•  struct sched_domain_topology_level {
•      sched_domain_mask_f mask;
•      sched_domain_flags_f sd_flags;
•      sched_domain_energy_f energy;
•      int flags;
•      int numa_level;
•      struct sd_data data;
• #ifdef CONFIG_SCHED_DEBUG
•      char *name;
• #endif
•  };

```

```

•
• struct sd_data {
•     struct sched_domain **__percpu sd;
•     struct sched_group **__percpu sg;
•     struct sched_group_capacity **__percpu sgc;
• };
•
• /*
•  * Topology list, bottom-up.
•  */
• static struct sched_domain_topology_level default_topology[] = {
• #ifdef CONFIG_SCHED_SMT
•     { cpu_smt_mask, cpu_smt_flags, SD_INIT_NAME(SMT) },
• #endif
• #ifdef CONFIG_SCHED_MC
•     { cpu_coregroup_mask, cpu_core_flags, SD_INIT_NAME(MC) },
• #endif
•     { cpu_cpu_mask, SD_INIT_NAME(DIE) },
•     { NULL, },
• };
•
• static struct sched_domain_topology_level *sched_domain_topology =
•     default_topology;

```

我们根据default\_topology[]这个结构体数组来解析sched\_domain\_topology\_level(简称SDTL)结构体成员变量的含义

- sched\_domain\_mask\_f mask:函数指针,用于指定某个SDTL层级的cpumask位图
- sched\_domain\_flags\_f sd\_flags:函数指针,用于指定某个SDTL层级的标志位
- sched\_domain\_energy\_f energy:函数指针,用于指定某个SDTL层级的energy信息,包括在active和idle状态下的capacity和power信息,具体获取是从dts(kernel/sched/energy.c)获取的([energy是怎么获取的](#))
- struct sd\_data:表示SDTL层级关系,包括domain/group/group capacity

从default\_topology结构体数组看,DIE是标配,MC和SMT是可选择的.目前ARM不支持SMT超线程技术,目的只有两种topology,MC和DIE.

由于在建立topology的时候,涉及到cpumask的使用,下面几个是比较常用的cpumask:

```

• #ifdef CONFIG_INIT_ALL_POSSIBLE
• static DECLARE_BITMAP(cpu_possible_bits, CONFIG_NR_CPUS) __read_mostly
•     = CPU_BITS_ALL;
• #else
• static DECLARE_BITMAP(cpu_possible_bits, CONFIG_NR_CPUS) __read_mostly;
• #endif
• const struct cpumask *const cpu_possible_mask =
•     to_cpumask(cpu_possible_bits);
• EXPORT_SYMBOL(cpu_possible_mask);
•
• static DECLARE_BITMAP(cpu_online_bits, CONFIG_NR_CPUS) __read_mostly;
• const struct cpumask *const cpu_online_mask = to_cpumask(cpu_online_bits);
• EXPORT_SYMBOL(cpu_online_mask);
•
• static DECLARE_BITMAP(cpu_present_bits, CONFIG_NR_CPUS) __read_mostly;
• const struct cpumask *const cpu_present_mask = to_cpumask(cpu_present_bits);
• EXPORT_SYMBOL(cpu_present_mask);

```



```

•
• static DECLARE_BITMAP(cpu_active_bits, CONFIG_NR_CPUS) __read_mostly;
• const struct cpumask *const cpu_active_mask = to_cpumask(cpu_active_bits);
• EXPORT_SYMBOL(cpu_active_mask);

```

上面的信息通过他们的名字就可以理解,需要注意的一点就是.online mask与active mask区别是:

- cpu\_online\_mask:表示当前系统online的cpu,包括idle+active cpu
- cpu\_active\_mask:表示当前系统online并且active的cpu

下面看看cpu topology怎么建立起来并经过初始化建立起sched domain和sched group的关系的.当boot cpu启动从cpu的时候,cpu topology就开始建立:

start\_kernel-->rest\_init-->kernel\_init-->kernel\_init\_freeable-->sched\_init\_smp-->init\_sched\_domains:

```

• /*
•  * Set up scheduler domains and groups. Callers must hold the hotplug lock.
•  * For now this just excludes isolated cpus, but could be used to
•  * exclude other special cases in the future.
•  */
• static int init_sched_domains(const struct cpumask *cpu_map)
• {
•     int err;
•     /*ARM没有定义*/
•     arch_update_cpu_topology();
•     ndoms_cur = 1;
•     /*分配一个cpumask结构体变量*/
•     doms_cur = alloc_sched_domains(ndoms_cur);
•     if (!doms_cur)
•         doms_cur = &fallback_doms;
•     /*doms_curr[0] = cpu_map & (~cpu_isolated_map),cpumask的与非的计算方式
•     cpu_isolated_map表示有独立的domain的cpu,当前在建立的cpu topology过程中需要
•     剔除cpu_isolated_map cpu*/
•     cpumask_andnot(doms_cur[0], cpu_map, cpu_isolated_map);
•     /*开始建立cpu topology*/
•     err = build_sched_domains(doms_cur[0], NULL);
•     /*注册sched domain的系统控制接口*/
•     register_sched_domain_sysctl();
•
•     return err;
• }
• /*
•  * arch_update_cpu_topology lets virtualized architectures update the
•  * cpu core maps. It is supposed to return 1 if the topology changed
•  * or 0 if it stayed the same.
•  */
• int __weak arch_update_cpu_topology(void)
• {
•     return 0;
• }
•
• cpumask_var_t *alloc_sched_domains(unsigned int ndoms)
• {
•     int i;
•     cpumask_var_t *doms;
•
•     doms = kmalloc(sizeof(*doms) * ndoms, GFP_KERNEL);
•     if (!doms)
•         return NULL;

```

```

•     for (i = 0; i < ndoms; i++) {
•         if (!alloc_cpumask_var(&doms[i], GFP_KERNEL)) {
•             free_sched_domains(doms, i);
•             return NULL;
•         }
•     }
•     return doms;
• }
•
• /*上面的cpu_map是active_cpu_mask,那么active cpumask在哪里设置的呢?*/
• /* Called by boot processor to activate the rest. */
• /*arm_dt_init_cpu_maps()函数会初始化若干个cpumask变量比如possible和present
• cpumask这个函数会将所有的cpu online起来,同时会等到online cpumask和active cpumask*/
• void __init smp_init(void)
• {
•     unsigned int cpu;
•
•     idle_threads_init();
•
•     /* FIXME: This should be done in userspace --RR */
•     for_each_present_cpu(cpu) {
•         if (num_online_cpus() >= setup_max_cpus)
•             break;
•         if (!cpu_online(cpu))
•             cpu_up(cpu);
•     }
•
•     /* Any cleanup work */
•     smp_announce();
•     smp_cpus_done(setup_max_cpus);
• }
•

```

接下来解析真正构建cpu topology结构的函数:build\_sched\_domains

```

• /*
•  * Build sched domains for a given set of cpus and attach the sched domains
•  * to the individual cpus
•  */
• static int build_sched_domains(const struct cpumask *cpu_map,
•                               struct sched_domain_attr *attr)
• {
•     enum s_alloc alloc_state;
•     struct sched_domain *sd;
•     struct s_data d;
•     int i, ret = -ENOMEM;
•
•     alloc_state = __visit_domain_allocation_hell(&d, cpu_map);
•     if (alloc_state != sa_rootdomain)
•         goto error;
•
•     /* Set up domains for cpus specified by the cpu_map. */
•     for_each_cpu(i, cpu_map) {
•         struct sched_domain_topology_level *tl;
•
•         sd = NULL;
•         for_each_sd_topology(tl) {
•
•             sd = build_sched_domain(tl, cpu_map, attr, sd, i);
•             if (tl == sched_domain_topology)
•                 *per_cpu_ptr(d.sd, i) = sd;
•

```

```

        if (tl->flags & SDTL_OVERLAP || sched_feat(FORCE_SD_OVERLAP))
            sd->flags |= SD_OVERLAP;
    }
}

/* Build the groups for the domains */
for_each_cpu(i, cpu_map) {
    for (sd = *per_cpu_ptr(d.sd, i); sd; sd = sd->parent) {
        sd->span_weight = cpumask_weight(sched_domain_span(sd));
        if (sd->flags & SD_OVERLAP) {
            if (build_overlap_sched_groups(sd, i))
                goto error;
        } else {
            if (build_sched_groups(sd, i))
                goto error;
        }
    }
}

/* Calculate CPU capacity for physical packages and nodes */
for (i = nr_cpumask_bits-1; i >= 0; i--) {
    struct sched_domain_topology_level *tl = sched_domain_topology;

    if (!cpumask_test_cpu(i, cpu_map))
        continue;

    for (sd = *per_cpu_ptr(d.sd, i); sd; sd = sd->parent, tl++) {
        init_sched_energy(i, sd, tl->energy);
        claim_allocations(i, sd);
        init_sched_groups_capacity(i, sd);
    }
}

/* Attach the domains */
rcu_read_lock();
for_each_cpu(i, cpu_map) {
    int max_cpu = READ_ONCE(d.rd->max_cap_orig_cpu);
    int min_cpu = READ_ONCE(d.rd->min_cap_orig_cpu);

    if ((max_cpu < 0) || (cpu_rq(i)->cpu_capacity_orig >
        cpu_rq(max_cpu)->cpu_capacity_orig))
        WRITE_ONCE(d.rd->max_cap_orig_cpu, i);

    if ((min_cpu < 0) || (cpu_rq(i)->cpu_capacity_orig <
        cpu_rq(min_cpu)->cpu_capacity_orig))
        WRITE_ONCE(d.rd->min_cap_orig_cpu, i);

    sd = *per_cpu_ptr(d.sd, i);

    cpu_attach_domain(sd, d.rd, i);
}
rcu_read_unlock();

ret = 0;

```

```

• error:
•     __free_domain_allocs(&d, alloc_state, cpu_map);
•     return ret;
• }

```

对构建cpu topology函数分如下几个部分来分析:

1.分配空间:在build\_sched\_domains函数里面调用\_\_visit\_domain\_allocation\_hell函数

```

• static enum s_alloc __visit_domain_allocation_hell(struct s_data *d,
•                                     const struct cpumask *cpu_map)
• {
•     memset(d, 0, sizeof(*d));
•     /*核心函数:创建调度域等数据结构,详细看下面的解析*/
•     if (__sdt_alloc(cpu_map))
•         return sa_sd_storage;
•     d->sd = alloc_percpu(struct sched_domain *);
•     if (!d->sd)
•         return sa_sd_storage;
•     d->rd = alloc_rootdomain();
•     if (!d->rd)
•         return sa_sd;
•     return sa_rootdomain;
• }
• static int __sdt_alloc(const struct cpumask *cpu_map)
• {
•     struct sched_domain_topology_level *tl;
•     int j;
•     /*
•     static struct sched_domain_topology_level *sched_domain_topology =
•         default_topology;
•     #define for_each_sd_topology(tl) \
•         for (tl = sched_domain_topology; tl->mask; tl++)
•     */ /*对default_topology结构体数据进行遍历*/
•     for_each_sd_topology(tl) {
•         /*获取每个SDTL层级的sd_data数据结构指针*/
•         struct sd_data *sdd = &tl->data;
•         /*下面三个alloc_percpu为每一个SDTL层级的实体数据结构sd_data分配空间
•         1.sched domain
•         2. sched group
•         3. sched group capacity*/
•         sdd->sd = alloc_percpu(struct sched_domain *);
•         if (!sdd->sd)
•             return -ENOMEM;
•
•         sdd->sg = alloc_percpu(struct sched_group *);
•         if (!sdd->sg)
•             return -ENOMEM;
•
•         sdd->sgc = alloc_percpu(struct sched_group_capacity *);
•         if (!sdd->sgc)
•             return -ENOMEM;
•
•         /*针对每个active_cpu_mask进行遍历,每个cpu上都需要为sd/sg/sgc分配空间,并存放在
•         percpu变量中.注意到sg->next指针,有两个不同的用途:
•         1.对于MC 层级,不同cluster的每个cpu core sched group会组成一个链表,比如两个
•         cluster,每个cluster 四个core,则每个cluster的sched group链表长度都为4,每个
•         core的sched domain都有一个sched group.

```

2. DIE层级,不同cluster的sched group会链接起来.一个cluster所有core的sched domain都共享一个sched group,后面会以图表说明\*/

```
for_each_cpu(j, cpu_map) {
    struct sched_domain *sd;
    struct sched_group *sg;
    struct sched_group_capacity *sgc;

    sd = kzalloc_node(sizeof(struct sched_domain) + cpumask_size(),
                      GFP_KERNEL, cpu_to_node(j));
    if (!sd)
        return -ENOMEM;

    *per_cpu_ptr(sdd->sd, j) = sd;

    sg = kzalloc_node(sizeof(struct sched_group) + cpumask_size(),
                      GFP_KERNEL, cpu_to_node(j));
    if (!sg)
        return -ENOMEM;

    sg->next = sg;

    *per_cpu_ptr(sdd->sg, j) = sg;

    sgc = kzalloc_node(sizeof(struct sched_group_capacity) +
                      cpumask_size(),
                      GFP_KERNEL, cpu_to_node(j));
    if (!sgc)
        return -ENOMEM;

    *per_cpu_ptr(sdd->sgc, j) = sgc;
}

return 0;
}static int __sdt_alloc(const struct cpumask *cpu_map)
{
    struct sched_domain_topology_level *tl;
    int j;

    for_each_sd_topology(tl) {
        struct sd_data *sdd = &tl->data;

        sdd->sd = alloc_percpu(struct sched_domain *);
        if (!sdd->sd)
            return -ENOMEM;

        sdd->sg = alloc_percpu(struct sched_group *);
        if (!sdd->sg)
            return -ENOMEM;

        sdd->sgc = alloc_percpu(struct sched_group_capacity *);
        if (!sdd->sgc)
            return -ENOMEM;
```

```

•     for_each_cpu(j, cpu_map) {
•         struct sched_domain *sd;
•         struct sched_group *sg;
•         struct sched_group_capacity *sgc;
•
•         sd = kzalloc_node(sizeof(struct sched_domain) + cpumask_size(),
•             GFP_KERNEL, cpu_to_node(j));
•         if (!sd)
•             return -ENOMEM;
•
•         *per_cpu_ptr(sdd->sd, j) = sd;
•
•         sg = kzalloc_node(sizeof(struct sched_group) + cpumask_size(),
•             GFP_KERNEL, cpu_to_node(j));
•         if (!sg)
•             return -ENOMEM;
•
•         sg->next = sg;
•
•         *per_cpu_ptr(sdd->sg, j) = sg;
•
•         sgc = kzalloc_node(sizeof(struct sched_group_capacity) +
•             cpumask_size(),
•             GFP_KERNEL, cpu_to_node(j));
•         if (!sgc)
•             return -ENOMEM;
•
•         *per_cpu_ptr(sdd->sgc, j) = sgc;
•     }
•
•     return 0;
• }

```

说明如下:

- 每个SDTL层级都有一个struct sched\_domain\_topology\_level数据结构来描述,并且内嵌了一个struct sd\_data数据结构,包含 sched\_domain,sched\_group,sched\_group\_capacity的二级指针
- 每个SDTL层级都分配一个percpu变量 sched\_domain,sched\_group,sched\_group\_capacity数据结构
- 在每个SDTL层级中为每个cpu都分配 sched\_domain,sched\_group,sched\_group\_capacity数据结构,即每个cpu在没给SDTL层级中都有对应的调度域和调度组.

2. 初始化sched\_domain:build\_sched\_domain函数,\_\_visit\_domain\_allocation\_hell函数已经为每个SDTL层级分配了sched\_domain,sched\_group/sched\_group\_capacity空间了,有看空间之后,就开始为cpu\_map建立调度域关系了:

```

•     build_sched_domains-->build_sched_domain
•         /* Set up domains for cpus specified by the cpu_map. */
•         /*对每个cpu_map进行遍历,并对每个CPU的SDTL层级进行遍历,相当于每个CPU都有一整套SDTL对
•         应的调度域,为每个CPU都初始化一整套SDTL对应的调度域和调度组,每个
•         CPU的每个SDTL都调用build_sched_domain函数来建立调度域和调度组*/
•         for_each_cpu(i, cpu_map) {

```

```

•      struct sched_domain_topology_level *tl;
•
•      sd = NULL;
•      for_each_sd_topology(tl) {
•          sd = build_sched_domain(tl, cpu_map, attr, sd, i);
•          if (tl == sched_domain_topology)
•              *per_cpu_ptr(d.sd, i) = sd;
•          if (tl->flags & SDTL_OVERLAP || sched_feat(FORCE_SD_OVERLAP))
•              sd->flags |= SD_OVERLAP;
•      }
•  }
•  |->
•  static struct sched_domain *
•  sd_init(struct sched_domain_topology_level *tl,
•          struct sched_domain *child, int cpu)
•  { /*获取cpu的sched_domain实体*/
•      struct sched_domain *sd = *per_cpu_ptr(tl->data.sd, cpu);
•      int sd_weight, sd_flags = 0;
•
•      #ifdef CONFIG_NUMA
•          /*
•           * Ugly hack to pass state to sd_numa_mask()...
•           */
•          sched_domains_curr_level = tl->numa_level;
•      #endif
•      /*通过default_topology[]结构体数据填充对应的参数,涉及到一些cpumask,这个上面已经
•      讲过了*/
•      sd_weight = cpumask_weight(tl->mask(cpu));
•
•      if (tl->sd_flags)
•          sd_flags = (*tl->sd_flags)();
•      if (WARN_ONCE(sd_flags & ~TOPOLOGY_SD_FLAGS,
•          "wrong sd_flags in topology description\n"))
•          sd_flags &= ~TOPOLOGY_SD_FLAGS;
•      /*填充sched_domain一些成员变量*/
•      *sd = (struct sched_domain){
•          .min_interval      = sd_weight,
•          .max_interval      = 2*sd_weight,
•          .busy_factor       = 32,
•          .imbalance_pct     = 125,
•
•          .cache_nice_tries  = 0,
•          .busy_idx          = 0,
•          .idle_idx          = 0,
•          .newidle_idx       = 0,
•          .wake_idx         = 0,
•          .forkexec_idx      = 0,
•
•          .flags              = 1*SD_LOAD_BALANCE
•                               | 1*SD_BALANCE_NEWIDLE
•                               | 1*SD_BALANCE_EXEC
•                               | 1*SD_BALANCE_FORK
•                               | 0*SD_BALANCE_WAKE
•                               | 1*SD_WAKE_AFFINE

```

```

•         | 0*SD_SHARE_CPUCAPACITY
•         | 0*SD_SHARE_PKG_RESOURCES
•         | 0*SD_SERIALIZE
•         | 0*SD_PREFER_SIBLING
•         | 0*SD_NUMA
• #ifdef CONFIG_INTEL_DWS
•         | 0*SD_INTEL_DWS
• #endif
•         | sd_flags
•         ,
•
•         .last_balance      = jiffies,
•         .balance_interval  = sd_weight,
•         .smt_gain          = 0,
•         .max_newidle_lb_cost = 0,
•         .next_decay_max_lb_cost = jiffies,
•         .child             = child,
• #ifdef CONFIG_INTEL_DWS
•         .dws_tf            = 0,
• #endif
• #ifdef CONFIG_SCHED_DEBUG
•         .name              = tl->name,
• #endif
•     };
•
•     /*
•      * Convert topological properties into behaviour.
•      */
•
•     if (sd->flags & SD_ASYM_CPUCAPACITY) {
•         struct sched_domain *t = sd;
•
•         for_each_lower_domain(t)
•             t->flags |= SD_BALANCE_WAKE;
•     }
•
•     if (sd->flags & SD_SHARE_CPUCAPACITY) {
•         sd->flags |= SD_PREFER_SIBLING;
•         sd->imbalance_pct = 110;
•         sd->smt_gain = 1178; /* ~15% */
•
•     } else if (sd->flags & SD_SHARE_PKG_RESOURCES) {
•         sd->imbalance_pct = 117;
•         sd->cache_nice_tries = 1;
•         sd->busy_idx = 2;
•
• #ifdef CONFIG_NUMA
•     } else if (sd->flags & SD_NUMA) {
•         sd->cache_nice_tries = 2;
•         sd->busy_idx = 3;
•         sd->idle_idx = 2;
•
•         sd->flags |= SD_SERIALIZE;

```



```

    if (sched_domains_numa_distance[tl->numa_level] > RECLAIM_DISTANCE)
    {
        sd->flags &= ~(SD_BALANCE_EXEC |
            SD_BALANCE_FORK |
            SD_WAKE_AFFINE);
    }

#endif
    } else {
        sd->flags |= SD_PREFER_SIBLING;
        sd->cache_nice_tries = 1;
        sd->busy_idx = 2;
        sd->idle_idx = 1;
    }

#ifdef CONFIG_INTEL_DWS
    if (sd->flags & SD_INTEL_DWS)
        sd->dws_tf = 120;
    if (sd->flags & SD_INTEL_DWS && sd->flags & SD_SHARE_PKG_RESOURCES)
        sd->dws_tf = 160;
#endif

    sd->private = &tl->data;

    return sd;
}

struct sched_domain *build_sched_domain(struct sched_domain_topology_level
*tl,
    const struct cpumask *cpu_map, struct sched_domain_attr *attr,
    struct sched_domain *child, int cpu)
{
    /*上面解释了sd_init函数*/
    struct sched_domain *sd = sd_init(tl, child, cpu);

    cpumask_and(sched_domain_span(sd), cpu_map, tl->mask(cpu));
    if (child) {
        sd->level = child->level + 1;
        sched_domain_level_max = max(sched_domain_level_max, sd->level);
        child->parent = sd;

        if (!cpumask_subset(sched_domain_span(child),
            sched_domain_span(sd))) {
            pr_err("BUG: arch topology borken\n");
#ifdef CONFIG_SCHED_DEBUG
            pr_err("    the %s domain not a subset of the %s domain\n",
                child->name, sd->name);
#endif
        }

        /* Fixup, ensure @sd has at least @child cpus. */
        cpumask_or(sched_domain_span(sd),
            sched_domain_span(sd),
            sched_domain_span(child));
    }
}

```

```

•     }
•     set_domain_attribute(sd, attr);
•
•     return sd;
• }

```

说明如下:

- sd\_init函数通过default\_topology结构体填充sched\_domain结构体并填充相关的一些参数
- cpumask\_and(sched\_domain\_span(sd), cpu\_map, tl->mask(cpu)); tl->mask(cpu)返回该cpu某个SDTL层级下兄弟cpu的bitmap, cpumask\_and则将tl->mask(cpu)的bitmap复制到span数组中.
- 我们知道遍历的顺序是SMT-->MC-->DIE, 可以看成父子关系或者上下级关系, struct sched\_domain数据结构中有parent和child成员用于描述此关系.
- 后面完成了对调度域的初始化.

### 3.初始化sched\_group.

```

• build_sched_domains--->build_sched_groups
•
• /* Build the groups for the domains */
• for_each_cpu(i, cpu_map) {
•     /*per_cpu_ptr(d.sd, i) 获取每个cpu的最低层级的SDTL, 并且通过指针sd->parent遍
•     历所有层级*/
•     for (sd = *per_cpu_ptr(d.sd, i); sd; sd = sd->parent) {
•         /*cpumask_and(sched_domain_span(sd), cpu_map, tl->mask(cpu))
•         某个SDTL层级cpu兄弟 cpu bitmap,*/
•         sd->span_weight = cpumask_weight(sched_domain_span(sd));
•         if (sd->flags & SD_OVERLAP) {
•             if (build_overlap_sched_groups(sd, i))
•                 goto error;
•         } else {
•             if (build_sched_groups(sd, i))
•                 goto error;
•         }
•     }
• }
•
• -->
• /*
•  * build_sched_groups will build a circular linked list of the groups
•  * covered by the given span, and will set each group's ->cpumask correctly,
•  * and ->cpu_capacity to 0.
•  *
•  * Assumes the sched_domain tree is fully constructed
•  * 核心函数
•  */
• static int
• build_sched_groups(struct sched_domain *sd, int cpu)
• {
•     struct sched_group *first = NULL, *last = NULL;
•     struct sd_data *sdd = sd->private;
•     const struct cpumask *span = sched_domain_span(sd);
•     struct cpumask *covered;
•     int i;
•
•

```

```

• get_group(cpu, sdd, &sd->groups);
• atomic_inc(&sd->groups->ref);
•
• if (cpu != cpumask_first(span))
•     return 0;
•
• lockdep_assert_held(&sched_domains_mutex);
• covered = sched_domains_tmpmask;
•
• cpumask_clear(covered);
• /*span是某个SDTL层级的cpu bitmap,实质某个sched domain包含多少个cpu*/
• /*两个for循环依次设置了该调度域sd中不同CPU对应的调度组的包含关系,这些调度组通过next
• 指针串联起来*/
• for_each_cpu(i, span) {
•     struct sched_group *sg;
•     int group, j;
•
•     if (cpumask_test_cpu(i, covered))
•         continue;
•
•     group = get_group(i, sdd, &sg);
•     cpumask_setall(sched_group_mask(sg));
•
•     for_each_cpu(j, span) {
•         if (get_group(j, sdd, NULL) != group)
•             continue;
•
•         cpumask_set_cpu(j, covered);
•         cpumask_set_cpu(j, sched_group_cpus(sg));
•     }
•
•     if (!first)
•         first = sg;
•     if (last)
•         last->next = sg;
•     last = sg;
• }
• /*将一个sched domain里面的所有sched group串成一组链表*/
• last->next = first;
•
• return 0;
• }
• static int get_group(int cpu, struct sd_data *sdd, struct sched_group **sg)
• {
•     struct sched_domain *sd = *per_cpu_ptr(sdd->sd, cpu);
•     struct sched_domain *child = sd->child;
•     /*如果是SDTL=DIE,那么child为true*/
•     if (child)
•         cpu = cpumask_first(sched_domain_span(child));
•
•     if (sg) {
•         *sg = *per_cpu_ptr(sdd->sg, cpu);
•         (*sg)->sgc = *per_cpu_ptr(sdd->sgc, cpu);
•         atomic_set(&(*sg)->sgc->ref, 1); /* for claim_allocations */
•     }
•
•     return cpu;

```

```

• }
•

```

#### 4 初始化sched group capacity:

```

• /* Calculate CPU capacity for physical packages and nodes */
• for (i = nr_cpumask_bits-1; i >= 0; i--) {
•     struct sched_domain_topology_level *tl = sched_domain_topology;
•
•     if (!cpumask_test_cpu(i, cpu_map))
•         continue;
•
•     for (sd = *per_cpu_ptr(d.sd, i); sd; sd = sd->parent, tl++) {
•         init_sched_energy(i, sd, tl->energy);
•         claim_allocations(i, sd);
•         init_sched_groups_capacity(i, sd);
•     }
• }
•
• /*上面的函数目的是填充下面的结构体每个cpu的每个SDTL层级的sched group capacity*/
• struct sched_group_capacity {
•     atomic_t ref;
•     /*
•      * CPU capacity of this group, SCHED_LOAD_SCALE being max capacity
•      * for a single CPU.
•      */
•     unsigned long capacity;
•     unsigned long max_capacity; /* Max per-cpu capacity in group */
•     unsigned long min_capacity; /* Min per-CPU capacity in group */
•     unsigned long next_update;
•     int imbalance; /* XXX unrelated to capacity but shared group state */
•     /*
•      * Number of busy cpus in this group.
•      */
•     atomic_t nr_busy_cpus;
•
•     unsigned long cpumask[0]; /* iteration mask */
• };
•

```

#### 5. 将每个cpu的调度域与每个cpu的struct rq root domain成员关联起来

```

• /* Attach the domains */
• rcu_read_lock();
• for_each_cpu(i, cpu_map) {
•     int max_cpu = READ_ONCE(d.rd->max_cap_orig_cpu);
•     int min_cpu = READ_ONCE(d.rd->min_cap_orig_cpu);
•
•     if ((max_cpu < 0) || (cpu_rq(i)->cpu_capacity_orig >
•         cpu_rq(max_cpu)->cpu_capacity_orig))
•         WRITE_ONCE(d.rd->max_cap_orig_cpu, i);
•
•     if ((min_cpu < 0) || (cpu_rq(i)->cpu_capacity_orig <
•         cpu_rq(min_cpu)->cpu_capacity_orig))
•         WRITE_ONCE(d.rd->min_cap_orig_cpu, i);
•
•     sd = *per_cpu_ptr(d.sd, i);
• }

```

```

•      /*将每个cpu的调度域关联到每个cpu的rq的root_domain上*/
•      cpu_attach_domain(sd, d.rd, i);
•  }
•  rcu_read_unlock();
•
•  /*
•   * Attach the domain 'sd' to 'cpu' as its base domain.
•   Callers must
•   * hold the hotplug lock.
•   */
•  static void
•  cpu_attach_domain(struct sched_domain *sd, struct root_domain
•  *rd, int cpu)
•  {
•      struct rq *rq = cpu_rq(cpu);
•      struct sched_domain *tmp;
•
•      /* Remove the sched domains which do not contribute to
•      scheduling. */
•      for (tmp = sd; tmp; ) {
•          struct sched_domain *parent = tmp->parent;
•          if (!parent)
•              break;
•
•          if (sd_parent_degenerate(tmp, parent)) {
•              tmp->parent = parent->parent;
•              if (parent->parent)
•                  parent->parent->child = tmp;
•              /*
•               * Transfer SD_PREFER_SIBLING down in case of a
•               * degenerate parent; the spans match for this
•               * so the property transfers.
•               */
•              if (parent->flags & SD_PREFER_SIBLING)
•                  tmp->flags |= SD_PREFER_SIBLING;
•              destroy_sched_domain(parent, cpu);
•          } else
•              tmp = tmp->parent;
•      }
•
•      if (sd && sd_degenerate(sd)) {
•          tmp = sd;
•          sd = sd->parent;
•          destroy_sched_domain(tmp, cpu);
•          if (sd)
•              sd->child = NULL;
•      }
•
•      sched_domain_debug(sd, cpu);
•
•      rq_attach_root(rq, rd);

```

- tmp = rq->sd;
- rcu\_assign\_pointer(rq->sd, sd);
- destroy\_sched\_domains(tmp, cpu);
- 
- update\_top\_cache\_domain(cpu);
- 
- #ifdef CONFIG\_INTEL\_DWS
- update\_dws\_domain(sd, cpu);
- #endif
- }

通过调度域,调度组的建立过程,可以通过下面的图来直观的说明SDTL不同层级对应关系:





