

我们经常看到很多driver会或者cpufreq governor等等会创建一些进程,并在创建之后使用wake_up_process(struct task_struct *p)函数直接wakeup这个task,下面看看这个wake_up_process函数是怎么实现进程调度,怎么实现放在哪个CPU上执行调度的?先列出起源码如下:

```
• wake_up_process(p) ----> try_to_wake_up(p, TASK_NORMAL, 0, 1) ---->
•
• /**
•  * wake_up_process - Wake up a specific process
•  * @p: The process to be woken up.
•  *
•  * Attempt to wake up the nominated process and move it to the set of
runnable
•  * processes.
•  *
•  * Return: 1 if the process was woken up, 0 if it was already running.
•  *
•  * It may be assumed that this function implies a write memory barrier
before
•  * changing the task state if and only if any tasks are woken up.
•  */
• int wake_up_process(struct task_struct *p)
• { /*wake_up_process直接调用try_to_wake_up函数,并添加三个限定参数*/
•     return try_to_wake_up(p, TASK_NORMAL, 0, 1);
• }
• /* Convenience macros for the sake of wake_up */
• #define TASK_NORMAL      (TASK_INTERRUPTIBLE | TASK_UNINTERRUPTIBLE)
•
• /**
•  * try_to_wake_up - wake up a thread
•  * @p: the thread to be awakened
•  * @state: the mask of task states that can be woken
•  * @wake_flags: wake modifier flags (WF_*)
•  * @sibling_count_hint: A hint at the number of threads that are being woken
up
•  *
•  *                               in this event.
•  *
•  * Put it on the run-queue if it's not already there. The "current"
•  * thread is always on the run-queue (except when the actual
•  * re-schedule is in progress), and as such you're allowed to do
•  * the simpler "current->state = TASK_RUNNING" to mark yourself
•  * runnable without the overhead of this.
•  *
•  * Return: %true if @p was woken up, %false if it was already running.
•  * or @state didn't match @p's state.
•  */
• static int
• try_to_wake_up(struct task_struct *p, unsigned int state, int wake_flags,
•               int sibling_count_hint)
• {
•     unsigned long flags;
•     int cpu, success = 0;
• #ifdef CONFIG_SMP
•     struct rq *rq;
•     u64 wallclock;
```

```

• #endif
•
• /*
• * If we are going to wake up a thread waiting for CONDITION we
• * need to ensure that CONDITION=1 done by the caller can not be
• * reordered with p->state check below. This pairs with mb() in
• * set_current_state() the waiting thread does.
• */ /*很有可能需要唤醒一个thread的函数,某个条件必须成立,为了取到最新的没有优化的
• 条件数值,使用内存屏障来实现.*/
• smp_mb__before_spinlock();
• raw_spin_lock_irqsave(&p->pi_lock, flags);
• /*如果进程不是在:TASK_INTERRUPTIBLE | TASK_UNINTERRUPTIBLE下,则就不是normal
• task,直接退出wakeup流程.所以在内核里面看到的wake_up_process,可以看到起主函数都会
• 将进程设置为TASK_INTERRUPTIBLE or TASK_UNINTERRUPTIBLE这两种状态之一*/
• if (!(p->state & state))
•     goto out;
•
• trace_sched_waking(p);
•
• success = 1; /* we're going to change ->state */
• /*获取这个进程当前处在的cpu上面,并不是时候进程就在这个cpu上运行,后面会挑选cpu*/
• cpu = task_cpu(p);
•
• /*
• * Ensure we load p->on_rq_after p->state, otherwise it would
• * be possible to, falsely, observe p->on_rq == 0 and get stuck
• * in smp_cond_load_acquire() below.
• *
• * sched_ttwa_pending()          try_to_wake_up()
• * [S] p->on_rq = 1;              [L] P->state
• *      UNLOCK rq->lock  -----.
• *                                \
• *      +--- RMB
• * schedule()                    /
• *      LOCK rq->lock  -----'
• *      UNLOCK rq->lock
• *
• * [task p]
• * [S] p->state = UNINTERRUPTIBLE [L] p->on_rq
• *
• * Pairs with the UNLOCK+LOCK on rq->lock from the
• * last wakeup of our task and the schedule that got our task
• * current.
• */
• smp_rmb();
• /*使用内存屏障保证p->on_rq的数值是最新的.如果task已经在rq里面,即进程已经处于
• runnable/running状态.ttwa_remote目的是由于task p已经在rq里面了,并且并没有完全
• 取消调度,再次会wakeup的话,需要将task的状态翻转,将状态设置为TASK_RUNNING,这样
• task就一直在rq里面运行了.这种情况直接退出下面的流程,并对调度状态/数据进行统计*/
• if (p->on_rq && ttwa_remote(p, wake_flags))
•     goto stat;
•
• #ifdef CONFIG_SMP
• /*

```

```

* Ensure we load p->on_cpu_after_ p->on_rq, otherwise it would be
* possible to, falsely, observe p->on_cpu == 0.
*
* One must be running (->on_cpu == 1) in order to remove oneself
* from the runqueue.
*
* [S] ->on_cpu = 1;    [L] ->on_rq
*     UNLOCK rq->lock
*     RMB
*     LOCK  rq->lock
* [S] ->on_rq = 0;    [L] ->on_cpu
*
* Pairs with the full barrier implied in the UNLOCK+LOCK on rq->lock
* from the consecutive calls to schedule(); the first switching to our
* task, the second putting it to sleep.
*/
smp_rmb();

/*
* If the owning (remote) cpu is still in the middle of schedule() with
* this task as prev, wait until its done referencing the task.
*/
/*如果拥有（远程）cpu仍然在schedule（）的中间，并且此任务为prev，请等待
其完成引用任务，意思是说，task p是作为其他cpu上的调度实体被调度，并且没有调度完毕，需要
等待其完毕，加内屏屏障就是保证on_cpu的数值是内存里面最新的数据
*/
while (p->on_cpu)
    cpu_relax(); /*cpu 忙等待，期望有人修改p->on_cpu的数值，退出忙等待*/

/*
* Combined with the control dependency above, we have an effective
* smp_load_acquire() without the need for full barriers.
*
* Pairs with the smp_store_release() in finish_lock_switch().
*
* This ensures that tasks getting woken will be fully ordered against
* their previous state and preserve Program Order.
*/
smp_rmb();
/*获取当前进程所在的cpu的rq*/
rq = cpu_rq(task_cpu(p));

raw_spin_lock(&rq->lock);
/*获取当前时间作为walt的wallclock*/
wallclock = walt_ktime_clock();
/*更新rq->curr在当前时间对进程负载/对累加的runnable_time的影响*/
walt_update_task_ravg(rq->curr, rq, TASK_UPDATE, wallclock, 0);
/*更新新创建的进程p在当前时间对进程负载/对累加的runnable_time的影响*/
walt_update_task_ravg(p, rq, TASK_WAKE, wallclock, 0);
raw_spin_unlock(&rq->lock);
/*根据进程p的状态来确定，如果调度器调度这个task是否会对load有贡献。*/
p->sched_contributes_to_load = !!task_contributes_to_load(p);
/*设置进程状态为TASK_WAKING*/
p->state = TASK_WAKING;

```

```

•    /*根据进程的所属的调度类调用相关的callback函数,这里进程会减去当前所处的cfs_rq的最小
•    vruntime,因为这里还没有确定进程会在哪个cpu上运行,等确定之后,会在入队的时候加上新cpu的
•    cfs_rq的最小vruntime*/
•    if (p->sched_class->task_waking)
•        p->sched_class->task_waking(p);
•    /*根据进程p相关参数设定和系统状态,为进程p选择合适的cpu供其运行*/
•    cpu = select_task_rq(p, p->wake_cpu, SD_BALANCE_WAKE, wake_flags,
•        sibling_count_hint);
•    /*如果选择的cpu与进程p当前所在的cpu不相同,则将进程的wake_flags标记添加需要迁移
•    ,并将进程p迁移到cpu上.*/
•    if (task_cpu(p) != cpu) {
•        wake_flags |= WF_MIGRATED;
•        set_task_cpu(p, cpu);
•    }
•
•    #endif /* CONFIG_SMP */
•    /*进程p入队操作并标记p为runnable状态,同时执行wakeup_preemption,即唤醒抢占.*/
•    ttwu_queue(p, cpu);
•    stat:
•        /*调度相关的统计*/
•        ttwu_stat(p, cpu, wake_flags);
•    out:
•        raw_spin_unlock_irqrestore(&p->pi_lock, flags);
•
•    return success;
• }
•

```

try_to_wake_up函数解析

try_to_wake_up函数的主要功能解析如下:

- 根据进程状态state来决定是否需要继续唤醒动作,这就是p->state & state的目的
- 获取进程p所在的当前cpu
- 如果当前cpu已经在rq里面了,则需要翻转进程p的状态为TASK_RUNNING,这样进程p就一直在rq里面,不需要继续调度,退出唤醒动作
- 如果进程p在cpu上,则期待某个时刻去修改这个数值on_cpu为0,这样代码可以继续运行.对于cpu_relax()真实含义如下:[cpu_relax\(\)真实含义](#)
- 根据WALT算法来实现rq当前task和新唤醒的task p相关的task load和rq相关的runnable_load的数值,具体怎么实现看:[WALT算法的原理分析](#)
- 为task p挑选合适的cpu
- 如果挑选的cpu与进程p所在的cpu不是同一个cpu,则进程task migration操作
- 将进程p入队操作
- 统计调度相关信息作为debug使用.

对于关键的几个点,详细分析如下:

1 How does the scheduler pick a suitable CPU?

```

•    cpu = select_task_rq(p, p->wake_cpu, SD_BALANCE_WAKE, wake_flags,
•        sibling_count_hint);

```

对函数select_task_rq的实现原理如下:

```

•    /*
•    * The caller (fork, wakeup) owns p->pi_lock, ->cpus_allowed is stable.
•    */
•    static inline

```

```

• int select_task_rq(struct task_struct *p, int cpu, int sd_flags, int
wake_flags,
•         int sibling_count_hint)
• {
•     lockdep_assert_held(&p->pi_lock);
•     /*nr_cpus_allowed这个变量是进程p可以运行的cpu数量,一般在系统初始化的时候就已经
•     设定好了的,或者可以通过设定cpu的亲 and 性来修改*/
•     if (p->nr_cpus_allowed > 1)
•         /*核心函数的callback*/
•         cpu = p->sched_class->select_task_rq(p, cpu, sd_flags, wake_flags,
•         sibling_count_hint);
•
•     /*
•     * In order not to call set_task_cpu() on a blocking task we need
•     * to rely on ttwu() to place the task on a valid ->cpus_allowed
•     * cpu.
•     *
•     * Since this is common to all placement strategies, this lives here.
•     *
•     * [ this allows ->select_task() to simply return task_cpu(p) and
•     *   not worry about this generic constraint ]
•     */
•     /*1. 如果选择的cpu与进程p允许运行的cpu不匹配 或者
•     2. 如果挑选的cpu offline
•     只要满足上面的任何一条,则重新选择cpu在进程p成员变cpus_allowed里面选择*/
•     if (unlikely(!cpumask_test_cpu(cpu, tsk_cpus_allowed(p)) ||
•         !cpu_online(cpu)))
•         cpu = select_fallback_rq(task_cpu(p), p);
•
•     return cpu;
• }

```

所以select_task_rq函数分两部分来分析:

1.1 select_task_rq callback函数分析

```

•         cpu = p->sched_class->select_task_rq(p, cpu, sd_flags, wake_flags,
•         sibling_count_hint);
•     ----->
•     /*
•     * select_task_rq_fair: Select target runqueue for the waking task in
•     domains
•     * that have the 'sd_flag' flag set. In practice, this is SD_BALANCE_WAKE,
•     * SD_BALANCE_FORK, or SD_BALANCE_EXEC.
•     *
•     * Balances load by selecting the idlest cpu in the idlest group, or under
•     * certain conditions an idle sibling cpu if the domain has SD_WAKE_AFFINE
•     set.
•     *
•     * Returns the target cpu number.
•     *
•     * preempt must be disabled.
•     */
•     static int
•     select_task_rq_fair(struct task_struct *p, int prev_cpu, int sd_flag, int
wake_flags,
•         int sibling_count_hint)

```

```

{
    struct sched_domain *tmp, *affine_sd = NULL, *sd = NULL;
    int cpu = smp_processor_id(); /*获取当前运行的cpu id*/
    int new_cpu = prev_cpu; /*将唤醒此进程p的cpu作为new_cpu*/
    int want_affine = 0;
    /*wake_falgs=0,so sync=0*/
    int sync = wake_flags & WF_SYNC;

#ifdef CONFIG_64BIT_ONLY_CPU
    struct cpumask tmpmask;

    if (find_packing_cpu(p, &new_cpu))
        return new_cpu;

    cpumask_andnot(&tmpmask, cpu_present_mask, &b64_only_cpu_mask);
    if (cpumask_test_cpu(cpu, &tmpmask)) {
        if (weighted_cpuload_32bit(cpu) >
            sysctl_sched_32bit_load_threshold &&
            !test_tsk_thread_flag(p, TIF_32BIT))
            return min_load_64bit_only_cpu();
    }
#endif

    /*sd_flag = SD_BALANCE_WAKE,是成立的,want_affine是一个核心变量包括三个部分数值的&&,后面会详细分析*/
    if (sd_flag & SD_BALANCE_WAKE) {
        record_wakee(p);
        want_affine = !wake_wide(p, sibling_count_hint) &&
            !wake_cap(p, cpu, prev_cpu) &&
            cpumask_test_cpu(cpu, &p->cpus_allowed);
    }

    /*如果系统使用EAS来决策的,则走这个分支,这种重点,也是新加的调度方案,根据cpu的能效和capacity来挑选cpu*/
    if (energy_aware())
        return select_energy_cpu_brute(p, prev_cpu, sync);

    rcu_read_lock();
    for_each_domain(cpu, tmp) {
        if (!(tmp->flags & SD_LOAD_BALANCE))
            break;

        /*
         * If both cpu and prev_cpu are part of this domain,
         * cpu is a valid SD_WAKE_AFFINE target.
         */
        if (want_affine && (tmp->flags & SD_WAKE_AFFINE) &&
            cpumask_test_cpu(prev_cpu, sched_domain_span(tmp))) {
            affine_sd = tmp;
            break;
        }
    }

    if (tmp->flags & sd_flag)
        sd = tmp;
    else if (!want_affine)
        break;
}

```

```

•     }
•
•     if (affine_sd) {
•         sd = NULL; /* Prefer wake_affine over balance flags */
•         if (cpu != prev_cpu && wake_affine(affine_sd, p, prev_cpu, sync))
•             new_cpu = cpu;
•     }
•
•     if (sd && !(sd_flag & SD_BALANCE_FORK)) {
•         /*
•          * We're going to need the task's util for capacity_spare_wake
•          * in find_idlest_group. Sync it up to prev_cpu's
•          * last_update_time.
•          */
•         sync_entity_load_avg(&p->se);
•     }
•
•     if (!sd) {
•         if (sd_flag & SD_BALANCE_WAKE) /* XXX always ? */
•             new_cpu = select_idle_sibling(p, prev_cpu, new_cpu);
•     } else {
•         new_cpu = find_idlest_cpu(sd, p, cpu, prev_cpu, sd_flag);
•     }
•     rcu_read_unlock();
•
•     return new_cpu;
• }

```

分别分几个部分来select_task_rq_fair函数:

1.1 want_affine变量怎么获取的呢?

```

•
•         want_affine = !wake_wide(p, sibling_count_hint) &&
•             !wake_cap(p, cpu, prev_cpu) &&
•             cpumask_test_cpu(cpu, &p->cpus_allowed);
•
•     ---->
•     /*
•      * Detect M:N waker/wakee relationships via a
•      * switching-frequency heuristic.
•      * A waker of many should wake a different task than the one
•      * last awakened
•      * at a frequency roughly N times higher than one of its
•      * wakees. In order
•      * to determine whether we should let the load spread vs
•      * consolodating to
•      * shared cache, we look for a minimum 'flip' frequency of
•      * llc_size in one
•      * partner, and a factor of lls_size higher frequency in the
•      * other. With
•      * both conditions met, we can be relatively sure that the
•      * relationship is
•      * non-monogamous, with partner count exceeding socket size.
•      * Waker/wakee

```

```

• * being client/server, worker/dispatcher, interrupt source
  or whatever is
• * irrelevant, spread criteria is apparent partner count
  exceeds socket size.
• */
• /*当前cpu的唤醒次数没有超标*/
• static int wake_wide(struct task_struct *p, int
  sibling_count_hint)
• {
•     unsigned int master = current->wakee_flips;
•     unsigned int slave = p->wakee_flips;
•     int llc_size = this_cpu_read(sd_llc_size);
•
•     if (sibling_count_hint >= llc_size)
•         return 1;
•
•     if (master < slave)
•         swap(master, slave);
•     if (slave < llc_size || master < slave * llc_size)
•         return 0;
•     return 1;
• }
•
• /*
•  * Disable WAKE_AFFINE in the case where task @p doesn't fit
  in the
•  * capacity of either the waking CPU @cpu or the previous CPU
  @prev_cpu.
•  *
•  * In that case WAKE_AFFINE doesn't make sense and we'll let
•  * BALANCE_WAKE sort things out.
•  */
• static int wake_cap(struct task_struct *p, int cpu, int
  prev_cpu)
• {
•     long min_cap, max_cap;
•     /*获取当前cpu的orig_of和唤醒进程p的cpu的orig_of capacity的最小
  值
•  */
•     min_cap = min(capacity_orig_of(prev_cpu),
  capacity_orig_of(cpu));
•     /*获取最大的capacity,为1024*/
•     max_cap = cpu_rq(cpu)->rd->max_cpu_capacity.val;
•
•     /* Minimum capacity is close to max, no need to abort
  wake_affine */
•     if (max_cap - min_cap < max_cap >> 3)
•         return 0;
•     /*根据PELT算法更新进程p作为调度实体的负载*/
•     /* Bring task utilization in sync with prev_cpu */
•     sync_entity_load_avg(&p->se);

```



```

    /*根据条件判断min_cap的capacity能够能够满足进程p吗?*/
    /*min_cap * 1024 < task_util(p) * 1138,
       task_util(p) ∈ [0,1024]*/
    return min_cap * 1024 < task_util(p) * capacity_margin;
}

static inline unsigned long task_util(struct task_struct *p)
{ /*WALT启用*/
#ifdef CONFIG_SCHED_WALT
    if (!walt_disabled && sysctl_sched_use_walt_task_util) {
        unsigned long demand = p->ravg.demand; /*task的真实运行时间*/
        /*在一个窗口内是多少,注意是*了1024的,比如占用了50%的窗口时间,则这个
        task_util = 0.5 * 1024=512.*/
        return (demand << 10) / walt_ravg_window;
    }
#endif
    return p->se.avg.util_avg;
}

/*
 * Synchronize entity load avg of dequeued entity without locking
 * the previous rq.
 */
void sync_entity_load_avg(struct sched_entity *se)
{
    struct cfs_rq *cfs_rq = cfs_rq_of(se);
    u64 last_update_time;

    last_update_time = cfs_rq_last_update_time(cfs_rq);
    /*PELT计算sched_entity调度实体的负载*/
    __update_load_avg(last_update_time, cpu_of(rq_of(cfs_rq)), &se->avg, 0, 0,
    NULL);
}

/**
 * cpumask_test_cpu - test for a cpu in a cpumask
 * @cpu: cpu number (< nr_cpu_ids)
 * @cpumask: the cpumask pointer
 *
 * Returns 1 if @cpu is set in @cpumask, else returns 0
 */
/*当前运行的cpu是否是task p cpu亲和数里面的一个*/
static inline int cpumask_test_cpu(int cpu, const struct
cpumask *cpumask)
{
    return test_bit(cpumask_check(cpu),
cpumask_bits((cpumask)));
}

```

只有满足下面三个条件:

- 当前cpu的唤醒次数没有超标
- 当前task p消耗的capacity * 1138小于min_cap * 1024
- 当前cpu在task p的cpu亲和数里面的一个

只有上面三个条件全部成立,则want_affine=1

1.2 使用EAS调度,怎么挑选合理的cpu呢?

如果使用EAS调度算法,则energy_aware()为true:

```
•     if (energy_aware())
•         return select_energy_cpu_brute(p, prev_cpu, sync);
•
•     static inline bool energy_aware(void)
•     { /*energy_aware调度类*/
•         return sched_feat(ENERGY_AWARE);
•     }
•
•     static int select_energy_cpu_brute(struct task_struct *p, int prev_cpu, int
sync)
•     {
•         struct sched_domain *sd;
•         int target_cpu = prev_cpu, tmp_target, tmp_backup;
•         bool boosted, prefer_idle;
•         /*调度统计信息*/
•         schedstat_inc(p, se.statistics.nr_wakeups_secb_attempts);
•         schedstat_inc(this_rq(), eas_stats.secb_attempts);
•         /*条件不成立*/
•         if (sysctl_sched_sync_hint_enable && sync) {
•             int cpu = smp_processor_id();
•
•             if (cpumask_test_cpu(cpu, tsk_cpus_allowed(p))) {
•                 schedstat_inc(p, se.statistics.nr_wakeups_secb_sync);
•                 schedstat_inc(this_rq(), eas_stats.secb_sync);
•                 return cpu;
•             }
•         }
•
•         rcu_read_lock();
•         /*下面的两个参数都可以在init.rc里面配置,一般boost都会配置,尤其是top-app*/
•         #ifdef CONFIG_CGROUP_SCHEDTUNE
•             /*获取当前task是否有util boost增益.如果有则boosted=true.
•             即如果原先的负载为util,那么boost之后的负载为util+boost/100*util*/
•             boosted = schedtune_task_boost(p) > 0;
•             /*获取在挑选cpu的时候,是否更倾向于idle cpu,默认为0,也就是说 prefer_idle=false*/
•             prefer_idle = schedtune_prefer_idle(p) > 0;
•         #else
•             boosted = get_sysctl_sched_cfs_boost() > 0;
•             prefer_idle = 0;
•         #endif
•         /*再次更新调度实体负载,使用PELT算法,比较奇怪的时候,在计算want_affine→
•         wake_cap函数里面已经update了调度实体的负载了,为何在这里还需要再次计算呢?*/
•         sync_entity_load_avg(&p->se);
•         /*DEFINE_PER_CPU(struct sched_domain *, sd_ea),在解析调度域调度组的创建和初始
•         化的时候,解析过,每个cpu在每个SDTL上面都有对应的调度域*/
•         sd = rcu_dereference(per_cpu(sd_ea, prev_cpu));
•         /* Find a cpu with sufficient capacity */
•         /*核心函数*/
•         tmp_target = find_best_target(p, &tmp_backup, boosted, prefer_idle);
•     }
```

1.2.1 下面讲解核心函数find_best_target,函数比较长:

```
1. static inline int find_best_target(struct task_struct *p, int *backup_cpu,
2.                                   bool boosted, bool prefer_idle)
3. {
4.     unsigned long best_idle_min_cap_orig = ULONG_MAX;
5.     /*计算task p经过boost之后的util数值,即在task_util(p)的基础上+boost%*util*/
6.     unsigned long min_util = boosted_task_util(p);
7.     unsigned long target_capacity = ULONG_MAX;
8.     unsigned long min_wake_util = ULONG_MAX;
9.     unsigned long target_max_spare_cap = 0;
10.    int best_idle_cstate = INT_MAX;
11.    unsigned long target_cap = ULONG_MAX;
12.    unsigned long best_idle_cap_orig = ULONG_MAX;
13.    int best_idle = INT_MAX;
14.    int backup_idle_cpu = -1;
15.    struct sched_domain *sd;
16.    struct sched_group *sg;
17.    int best_active_cpu = -1;
18.    int best_idle_cpu = -1;
19.    int target_cpu = -1;
20.    int cpu, i;
21.    /*获取当前cpu的运行队列rq的root_domain*/
22.    struct root_domain *rd = cpu_rq(smp_processor_id())->rd;
23.    /*获取当前root_domain的最大capacity数值*/
24.    unsigned long max_cap = rd->max_cpu_capacity.val;
25.
26.    *backup_cpu = -1;
27.
28.    schedstat_inc(p, se.statistics.nr_wakeups_fbt_attempts);
29.    schedstat_inc(this_rq(), eas_stats.fbt_attempts);
30.
31.    /* Find start CPU based on boost value */
32.    /*start_cpu找出rd->min_cap_orig_cpu,即min_cap_orig的第一个cpu id
33.    min的cpu为0*/
34.    cpu = start_cpu(boosted);
35.    if (cpu < 0) {
36.        schedstat_inc(p, se.statistics.nr_wakeups_fbt_no_cpu);
37.        schedstat_inc(this_rq(), eas_stats.fbt_no_cpu);
38.        return -1;
39.    }
40.
41.    /* Find SD for the start CPU */
42.    /*找到启动cpu的调度域*/
43.    sd = rcu_dereference(per_cpu(sd_ea, cpu));
44.    if (!sd) {
45.        schedstat_inc(p, se.statistics.nr_wakeups_fbt_no_sd);
46.        schedstat_inc(this_rq(), eas_stats.fbt_no_sd);
47.        return -1;
48.    }
49.
50.    /* Scan CPUs in all SDs */
51.    sg = sd->groups;
52.    do {
```

```

53.     for_each_cpu_and(i, tsk_cpus_allowed(p), sched_group_cpus(sg)) {
54.         unsigned long capacity_orig = capacity_orig_of(i);
55.         unsigned long wake_util, new_util;
56.
57.         if (!cpu_online(i))
58.             continue;
59.
60.         if (walt_cpu_high_irqload(i))
61.             continue;
62.
63.         /*
64.          * p's blocked utilization is still accounted for on prev_cpu
65.          * so prev_cpu will receive a negative bias due to the double
66.          * accounting. However, the blocked utilization may be zero.
67.          */
68.         wake_util = cpu_util_wake(i, p);
69.         new_util = wake_util + task_util(p);
70.
71.         /*
72.          * Ensure minimum capacity to grant the required boost.
73.          * The target CPU can be already at a capacity level higher
74.          * than the one required to boost the task.
75.          */
76.         new_util = max(min_util, new_util);
77.         if (new_util > capacity_orig) {
78.             if (idle_cpu(i)) {
79.                 int idle_idx;
80.
81.                 idle_idx =
82.                     idle_get_state_idx(cpu_rq(i));
83.
84.                 if (capacity_orig >
85.                     best_idle_cap_orig) {
86.                     best_idle_cap_orig =
87.                         capacity_orig;
88.                     best_idle = idle_idx;
89.                     backup_idle_cpu = i;
90.                     continue;
91.                 }
92.
93.                 /*
94.                  * Skip CPUs in deeper idle state, but
95.                  * only if they are also less energy
96.                  * efficient.
97.                  * IOW, prefer a deep IDLE LITTLE CPU
98.                  * vs a shallow idle big CPU.
99.                  */
100.                if (sysctl_sched_cstate_aware &&
101.                    best_idle <= idle_idx)
102.                    continue;
103.
104.                /* Keep track of best idle CPU */
105.                best_idle_cap_orig = capacity_orig;
106.                best_idle = idle_idx;

```

```

107.             backup_idle_cpu = i;
108.             continue;
109.         }
110.
111.         if (capacity_orig > target_cap) {
112.             target_cap = capacity_orig;
113.             min_wake_util = wake_util;
114.             best_active_cpu = i;
115.             continue;
116.         }
117.
118.         if (wake_util > min_wake_util)
119.             continue;
120.
121.         min_wake_util = wake_util;
122.         best_active_cpu = i;
123.         continue;
124.
125.     }
126.     /*
127.      * Enforce EAS mode
128.      *
129.      * For non latency sensitive tasks, skip CPUs that
130.      * will be overutilized by moving the task there.
131.      *
132.      * The goal here is to remain in EAS mode as long as
133.      * possible at least for !prefer_idle tasks.
134.      */
135.     if (capacity_orig == max_cap)
136.         if (idle_cpu(i))
137.             goto skip;
138.
139.     if ((new_util * capacity_margin) >
140.         (capacity_orig * SCHED_CAPACITY_SCALE))
141.         continue;
142. skip:
143.     if (idle_cpu(i)) {
144.         int idle_idx;
145.
146.         if (prefer_idle ||
147.             cpumask_test_cpu(i, &min_cap_cpu_mask)) {
148.             trace_sched_find_best_target(p,
149.                 prefer_idle, min_util, cpu,
150.                 best_idle_cpu, best_active_cpu,
151.                 i);
152.             return i;
153.         }
154.         idle_idx = idle_get_state_idx(cpu_rq(i));
155.
156.         /* Select idle CPU with lower cap_orig */
157.         if (capacity_orig > best_idle_min_cap_orig)
158.             continue;
159.
160.         /*

```

```

161.             * Skip CPUs in deeper idle state, but only
162.             * if they are also less energy efficient.
163.             * IOW, prefer a deep IDLE LITTLE CPU vs a
164.             * shallow idle big CPU.
165.             */
166.             if (sysctl_sched_cstate_aware &&
167.                 best_idle_cstate <= idle_idx)
168.                 continue;
169.
170.             /* Keep track of best idle CPU */
171.             best_idle_min_cap_orig = capacity_orig;
172.             best_idle_cstate = idle_idx;
173.             best_idle_cpu = i;
174.             continue;
175.         }
176.
177.         /* Favor CPUs with smaller capacity */
178.         if (capacity_orig > target_capacity)
179.             continue;
180.
181.         /* Favor CPUs with maximum spare capacity */
182.         if ((capacity_orig - new_util) < target_max_spare_cap)
183.             continue;
184.
185.         target_max_spare_cap = capacity_orig - new_util;
186.         target_capacity = capacity_orig;
187.         target_cpu = i;
188.     }
189.
190. } while (sg = sg->next, sg != sd->groups);
191.
192. /*
193.  * For non latency sensitive tasks, cases B and C in the previous
194.  * loop,
195.  * we pick the best IDLE CPU only if we was not able to find a target
196.  * ACTIVE CPU.
197.  * Policies priorities:
198.  *
199.  * - prefer_idle tasks:
200.  *
201.  *   a) IDLE CPU available, we return immediately
202.  *   b) ACTIVE CPU where task fits and has the bigger maximum spare
203.  *       capacity (i.e. target_cpu)
204.  *   c) ACTIVE CPU with less contention due to other tasks
205.  *       (i.e. best_active_cpu)
206.  *
207.  * - NON prefer_idle tasks:
208.  *
209.  *   a) ACTIVE CPU: target_cpu
210.  *   b) IDLE CPU: best_idle_cpu
211.  */
212. if (target_cpu == -1) {
213.     if (best_idle_cpu != -1)

```

```

214.         target_cpu = best_idle_cpu;
215.     else
216.         target_cpu = (backup_idle_cpu != -1)
217.             ? backup_idle_cpu
218.             : best_active_cpu;
219.     } else
220.         *backup_cpu = best_idle_cpu;
221.
222.     trace_sched_find_best_target(p, prefer_idle, min_util, cpu,
223.         best_idle_cpu, best_active_cpu,
224.         target_cpu);
225.
226.     schedstat_inc(p, se.statistics.nr_wakeups_fbt_count);
227.     schedstat_inc(this_rq(), eas_stats.fbt_count);
228.
229.     return target_cpu;
230. }

```

分下面如下几个部分来分析上面的do{}while()循环

- do{}while()循环是对sched domain里面的所有调度组进行遍历
- for_each_cpu_and(i, tsk_cpus_allowed(p), sched_group_cpus(sg)),这个for循环比较有意思,sched_group_cpus(sg),表示这个调度组所有的cpumask,其返回值就是sg->cpumask.for_each_cpu_and抽象循环的意思是i必须在tsk_cpus_allowed(p) && sched_group_cpus(sg)交集里面。

```

/**
 * for_each_cpu_and - iterate over every cpu in both masks
 * @cpu: the (optionally unsigned) integer iterator
 * @mask: the first cpumask pointer
 * @and: the second cpumask pointer
 *
 * This saves a temporary CPU mask in many places. It is equivalent to:
 * struct cpumask tmp;
 * cpumask_and(&tmp, &mask, &and);
 * for_each_cpu(cpu, &tmp)
 * ...
 *
 * After the loop, cpu is >= nr_cpu_ids.
 */
#define for_each_cpu_and(cpu, mask, and) \
    for ((cpu) = -1; \
         (cpu) = cpumask_next_and((cpu), (mask), (and)), \
         (cpu) < nr_cpu_ids;)

```

- 如果cpu id为i的cpu offline,则遍历下一个cpu
- 如果这个cpu是一个irq high load的cpu,则遍历下一个cpu
- 计算wake_util,即为当前cpu id=i的cpu_util的数值,new_util为cpu i的cpu_util+进程p的task_util数值,最后new_util = max(min_util, new_util);min_util数值是task_util boost之后的数值,如果没有boost,则min_util=task_util。

上面条件判断之后,并且获取了当前遍历cpu的util和新唤醒的进程task_util叠加到cpu_util变成new_util之后,通过capacity/util的比较来获取target_cpu,下面分析代码77~187行代码:

在遍历cpu的时候

如果new_util > 遍历的cpu的capacity数值(dts获取),分两部分逻辑处理:

1. 如果cpu是idle状态.记录此cpu处在idle的level_idx,并修改下面三个参数数值之后遍历下一个符合条件的cpu:

- /*获取当前遍历cpu的capacity,并保存在best_idle_cap_orig变量中 */
- best_idle_cap_orig = capacity_orig;
- best_idle = idle_idx; //获取idle level index
- backup_idle_cpu = i; //idle cpu number,后面会使用到

2.如果不是idle cpu,则修正下面两个参数之后遍历下一个符合条件的cpu:

- min_wake_util = wake_util; //将遍历的cpu_util赋值给min_wake_util
- best_active_cpu = i; //得到best_active_cpu id

如果new_util <= 遍历的cpu的capacity数值(dts获取),分如下几部分逻辑处理:

1.如果capacity_orig == max_cap并且遍历的cpu恰好是idle状态,直接调到去更新下面三个参数:

- /*将当前遍历的cpu的capacity保存到best_idle_min_cap_orig变量中*/
- best_idle_min_cap_orig = capacity_orig;
- best_idle_cstate = idle_idx; //保存idle index
- best_idle_cpu = i; //保存最佳的idle cpu,后面会用到

2.如果(new_util * capacity_margin) > (capacity_orig * SCHED_CAPACITY_SCALE)成立,则直接遍历下一个cpu,说明此时遍历的cpu已经overutilization,没必要继续遍历了.

3.如果capacity_orig==max_cap不成立,也会执行第一条先判断遍历的cpu是否是idle,是的话,执行1一样的流程

4.如果遍历的cpu不是idle,则比较capacity_orig-new_util差值与target_max_spare_cap的比较目的是选择一个差值余量最大的cpu,防止新唤醒的task p在余量不足的cpu上运行导致后面的负载均衡,白白浪费系统资源.同时更新下面三个参数,并在每次遍历的时候更新:

- /*util余量,目的找出最大余量的cpu id*/
- target_max_spare_cap = capacity_orig - new_util;
- /*目标capacity*/
- target_capacity = capacity_orig;
- /*选择的目标cpu*/
- target_cpu = i;
-

下面分析212~220之间的代码:

从解释来看(对于非敏感延迟性进程)进程分两种,prefer_idle flag:

1. 一种是偏爱idle cpu运行的进程,那么如果有idle cpu,则优先选择idle cpu并立即返回,如代码145~152行的代码;之后task的util不太大,就选择有最大余量的cpu了;最后挑选有更少争抢的cpu,比如best_active_cpu
2. 不偏爱idle cpu运行的进程,优先选择余量最大的cpu,之后选择best_idle_cpu

明白了prefer_idle这个flag会影响对cpu类型的选择,那么分析212~220行之间的代码如下,即在有机会执行寻找最大余量的cpu capacity的情况下:

如果target_cpu = -1

- 如果best_idle_cpu更新过,表明要么new_util很大,要么大部分cpu处于idle状态,这时候直接选择best_idle_cpu为target_cpu
- 否则,根据backup_idle_cpu是否update过来决定target_cpu选择是backup_idle_cpu还是best_active_cpu

如果target_cpu != -1

- 候选cpu设置为best_idle_cpu,通过函数指针被使用

最后返回target_cpu的作为选择的cpu id.

1.2.2 select_energy_cpu_brute函数剩余部分分析:

```
static int select_energy_cpu_brute(struct task_struct *p, int prev_cpu, int
sync)
{
    .....
    /* Find a cpu with sufficient capacity */
    /*下面这个函数已经解析完毕*/
    tmp_target = find_best_target(p, &tmp_backup, boosted, prefer_idle);

    if (!sd)
        goto unlock;
    if (tmp_target >= 0) {
        target_cpu = tmp_target;
        /*如果boosted or prefer_idle && target_cpu为idlecpu,或者target_cpu为
        min_cap的cpu,则挑选的cpu就是target_cpu了并直接退出代码流程*/
        if (((boosted || prefer_idle) && idle_cpu(target_cpu)) ||
            cpumask_test_cpu(target_cpu, &min_cap_cpu_mask)) {
            schedstat_inc(p, se.statistics.nr_wakeups_secb_idle_bt);
            schedstat_inc(this_rq(), eas_stats.secb_idle_bt);
            goto unlock;
        }
    }
    /*如果target_cpu等于唤醒进程p的cpu并且best_idle_cpu>=0,则修改target_cpu为
    best_idle_cpu数值,目的不在唤醒进程P的cpu上运行.why??*/
    if (target_cpu == prev_cpu && tmp_backup >= 0) {
        target_cpu = tmp_backup;
        tmp_backup = -1;
    }

    if (target_cpu != prev_cpu) {
        int delta = 0;
        /*构造需要迁移的环境变量*/
        struct energy_env eenv = {
            .util_delta = task_util(p),
            .src_cpu = prev_cpu,
            .dst_cpu = target_cpu,
            .task = p,
            .trg_cpu = target_cpu,
        };

#ifdef CONFIG_SCHED_WALT
        if (!walt_disabled && sysctl_sched_use_walt_cpu_util &&
            p->state == TASK_WAKING)
            /*获取进程p本身的util load*/
            delta = task_util(p);
#endif

        /* Not enough spare capacity on previous cpu */
        /*唤醒进程p的cpu负载过载,超过本身capacity的90%.*/
        if (__cpu_overutilized(prev_cpu, delta, p)) {
            /*有限选择Energy合理的cpu,即小的cluster的idle cpu*/
            if (tmp_backup >= 0 &&
                capacity_orig_of(tmp_backup) <
                capacity_orig_of(target_cpu))
```

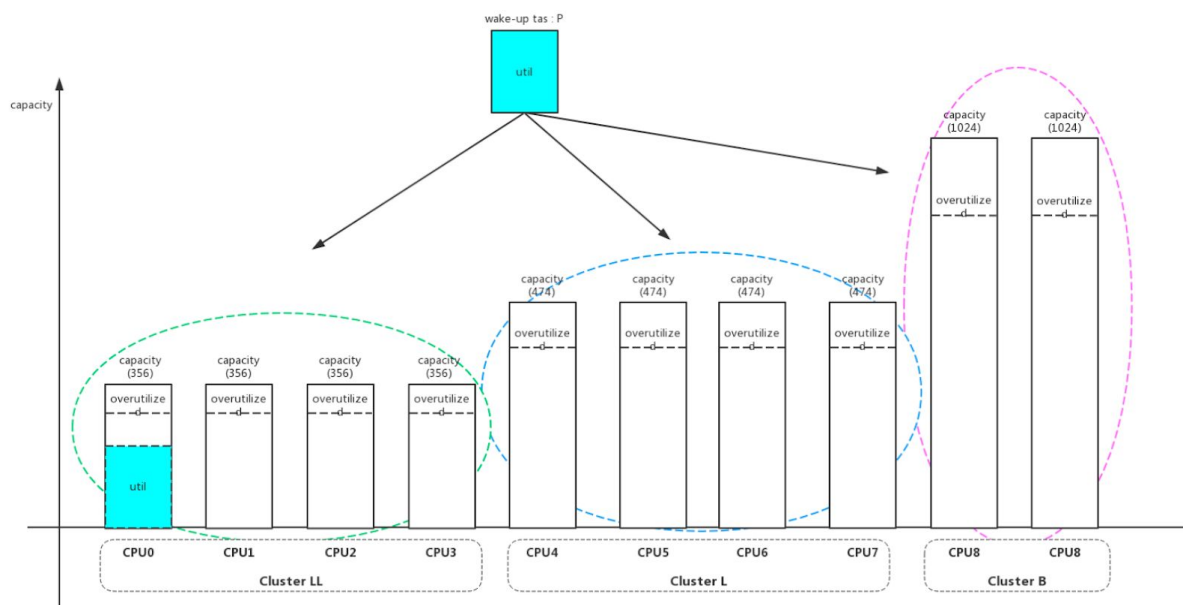
```

•         target_cpu = tmp_backup;
•         schedstat_inc(p, se.statistics.nr_wakeups_secb_insuff_cap);
•         schedstat_inc(this_rq(), eas_stats.secb_insuff_cap);
•         goto unlock;
•     }
•     /*计算pre_cpu与target_cpu的功耗差异, 如果大于0, 则执行下面的代码流程
•     就是计算MC知道DIE的SDTL的功耗总和的差异*/
•     if (energy_diff(&eenv) >= 0) {
•         /* No energy saving for target_cpu, try backup */
•         target_cpu = tmp_backup;
•         eenv.dst_cpu = target_cpu;
•         eenv.trg_cpu = target_cpu;
•         if (tmp_backup < 0 ||
•             tmp_backup == prev_cpu ||
•             energy_diff(&eenv) >= 0) {
•             schedstat_inc(p, se.statistics.nr_wakeups_secb_no_nrg_sav);
•             schedstat_inc(this_rq(), eas_stats.secb_no_nrg_sav);
•             target_cpu = prev_cpu;
•             goto unlock;
•         }
•     }
•
•     schedstat_inc(p, se.statistics.nr_wakeups_secb_nrg_sav);
•     schedstat_inc(this_rq(), eas_stats.secb_nrg_sav);
•     goto unlock;
• }
•
•     schedstat_inc(p, se.statistics.nr_wakeups_secb_count);
•     schedstat_inc(this_rq(), eas_stats.secb_count);
•
• unlock:
•     rcu_read_unlock();
•
•     return target_cpu;
• }

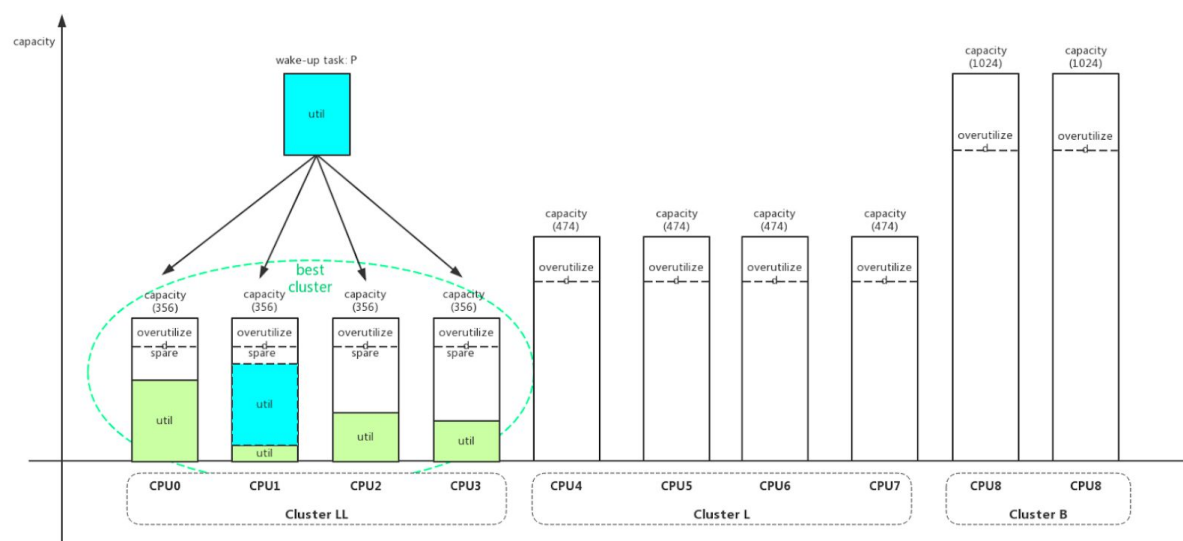
```

总结下就是如下四点:

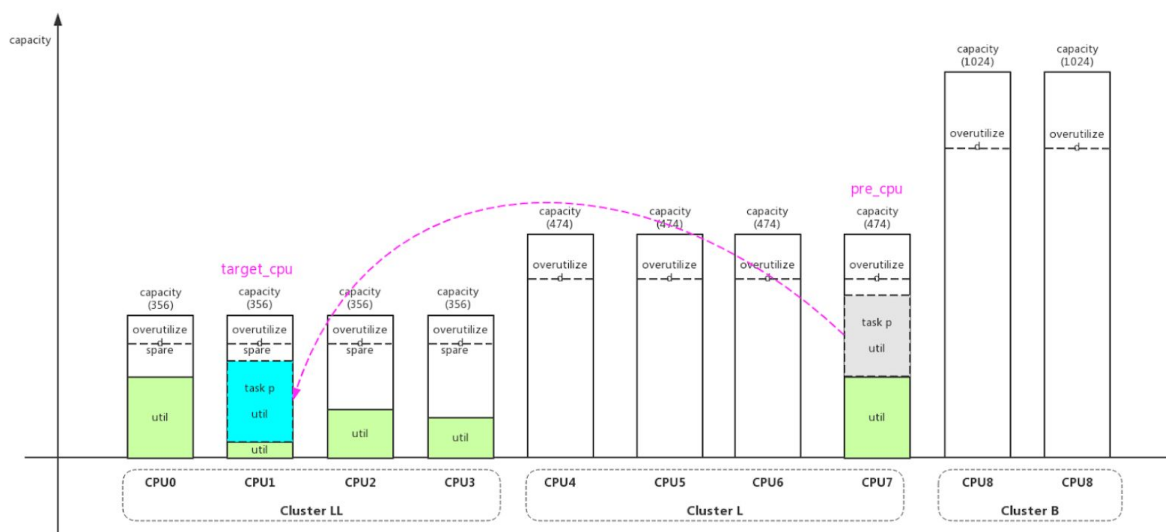
1. EAS遍历sched_group(cluster)和cpu, 找到一个既能满足进程p的affinity又能容纳下进程p的负载util, 属于能用最小capacity满足的cluster其中余量capacity最多的target_cpu. 首先找到能容纳进程p的util且capacity最小的cluster



2. 然后在目标cluster中找到加上进程p以后，剩余capacity最大的cpu

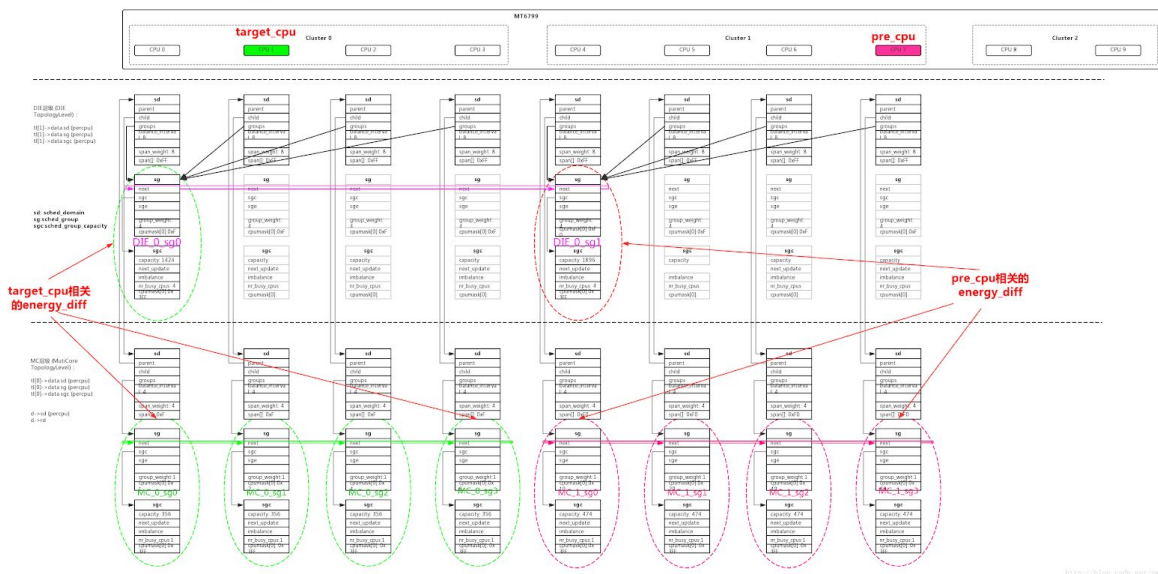


3. pre_cpu是进程p上一次运行的cpu作为src_cpu，上面选择的target_cpu作为dst_cpu，就是尝试计算进程p从pre_cpu迁移到target_cpu系统的功耗差异



4. 计算负载变化前后, target_cpu和prev_cpu带来的power变化。如果没有power增加则返回target_cpu, 如果有power增加则返回prev_cpu. 计算负载变化的函数energy_diff()循环很多比较复杂, 仔细分析下来就是计算target_cpu/prev_cpu在“MC层次cpu所在sg链表”+“DIE层级cpu所在sg”, 这两种范围在负载变化中的功耗差异:

energy_aware_wake_cpu_step4: energy_diff()



上面四个过程很清晰明了的pre_cpu和target_cpu的转换关系

接下来分析energy_diff(&env)怎么来计算pre_cpu与target_cpu power之间的关系。

```
static inline int
energy_diff(struct energy_env *eenv)
{
    int boost = schedtune_task_boost(eenv->task);
    int nrg delta;
    /* 计算绝对功耗差值 */
    /* Compute "absolute" energy diff */
    __energy_diff(eenv);
}
```

```

/* Return energy diff when boost margin is 0 */
if (1 || boost == 0) {
    trace_sched_energy_diff(eenv->task,
        eenv->src_cpu, eenv->dst_cpu, eenv->util_delta,
        eenv->nrg.before, eenv->nrg.after, eenv->nrg.diff,
        eenv->cap.before, eenv->cap.after, eenv->cap.delta,
        0, -eenv->nrg.diff);
    return eenv->nrg.diff;
}

/* Compute normalized energy diff */
nrg_delta = normalize_energy(eenv->nrg.diff);
eenv->nrg.delta = nrg_delta;

eenv->payoff = schedtune_accept_deltas(
    eenv->nrg.delta,
    eenv->cap.delta,
    eenv->task);

trace_sched_energy_diff(eenv->task,
    eenv->src_cpu, eenv->dst_cpu, eenv->util_delta,
    eenv->nrg.before, eenv->nrg.after, eenv->nrg.diff,
    eenv->cap.before, eenv->cap.after, eenv->cap.delta,
    eenv->nrg.delta, eenv->payoff);

/*
 * When SchedTune is enabled, the energy_diff() function will return
 * the computed energy payoff value. Since the energy_diff() return
 * value is expected to be negative by its callers, this evaluation
 * function return a negative value each time the evaluation return a
 * positive payoff, which is the condition for the acceptance of
 * a scheduling decision
 */
return -eenv->payoff;
}

/*
 * energy_diff(): Estimate the energy impact of changing the utilization
 * distribution. eenv specifies the change: utilisation amount, source, and
 * destination cpu. Source or destination cpu may be -1 in which case the
 * utilization is removed from or added to the system (e.g. task wake-up).
 */
If
/* both are specified, the utilization is migrated.
 */
static inline int __energy_diff(struct energy_env *eenv)
{
    struct sched_domain *sd;
    struct sched_group *sg;
    int sd_cpu = -1, energy_before = 0, energy_after = 0;
    int diff, margin;
    /*在brute函数里面设置好的eenv参数,构造迁移前的环境变量*/
    struct energy_env eenv_before = {
        .util_delta = task_util(eenv->task),
        .src_cpu = eenv->src_cpu,
        .dst_cpu = eenv->dst_cpu,
        .trg_cpu = eenv->src_cpu,
        .nrg = { 0, 0, 0, 0 },
        .cap = { 0, 0, 0 },
        .task = eenv->task,
    };

    if (eenv->src_cpu == eenv->dst_cpu)
        return 0;
    /*sd来至于cache sd_ea, 是cpu对应的顶层sd(tl DIE层)*/
    sd_cpu = (eenv->src_cpu != -1) ? eenv->src_cpu : eenv->dst_cpu;
    sd = rcu_dereference(per_cpu(sd_ea, sd_cpu));

```

```

if (!sd)
    return 0; /* Error */

sg = sd->groups;
/*遍历sg所在sg链表，找到符合条件的sg，累加计算eenv_before、eenv
相关sg的功耗*/
do { /*如果当前sg包含src_cpu或者dst_cpu，则进行计算*/
    if (cpu_in_sg(sg, eenv->src_cpu) || cpu_in_sg(sg, eenv->dst_cpu)) {
        /*当前顶层sg为eenv的sg_top*/
        eenv_before.sg_top = eenv->sg_top = sg;
        /*计算eenv_before负载下sg的power*/
        if (sched_group_energy(&eenv_before))
            return 0; /* Invalid result abort */
        energy_before += eenv_before.energy;

        /* Keep track of SRC cpu (before) capacity */
        eenv->cap.before = eenv_before.cap.before;
        eenv->cap.delta = eenv_before.cap.delta;
        /*计算eenv负载下sg的power*/
        if (sched_group_energy(eenv))
            return 0; /* Invalid result abort */
        energy_after += eenv->energy;
    }
} while (sg = sg->next, sg != sd->groups);
/*计算energy_after - energy_before*/
eenv->nrg.before = energy_before;
eenv->nrg.after = energy_after;
eenv->nrg.diff = eenv->nrg.after - eenv->nrg.before;
eenv->payoff = 0;
#ifdef CONFIG_SCHED_TUNE
    trace_sched_energy_diff(eenv->task,
        eenv->src_cpu, eenv->dst_cpu, eenv->util_delta,
        eenv->nrg.before, eenv->nrg.after, eenv->nrg.diff,
        eenv->cap.before, eenv->cap.after, eenv->cap.delta,
        eenv->nrg.delta, eenv->payoff);
#endif
/*
 * Dead-zone margin preventing too many migrations.
 */

margin = eenv->nrg.before >> 6; /* ~1.56% */

diff = eenv->nrg.after - eenv->nrg.before;

eenv->nrg.diff = (abs(diff) < margin) ? 0 : eenv->nrg.diff;

return eenv->nrg.diff;
}
/*接下来看sched_group_energy函数的实现过程*/
/*
 * sched_group_energy(): Computes the absolute energy consumption of cpus
 * belonging to the sched_group including shared resources shared only by
 * members of the group. Iterates over all cpus in the hierarchy below the
 * sched_group starting from the bottom working it's way up before going to
 * the next cpu until all cpus are covered at all levels. The current
 * implementation is likely to gather the same util statistics multiple
 * times.
 * This can probably be done in a faster but more complex way.
 * Note: sched_group_energy() may fail when racing with sched_domain
 * updates.
 */

```

```

static int sched_group_energy(struct energy_env *eenv)
{
    struct cpumask visit_cpus;
    u64 total_energy = 0;
    int cpu_count;

    WARN_ON(!eenv->sg_top->sge);

    cpumask_copy(&visit_cpus, sched_group_cpus(eenv->sg_top));
    /* If a cpu is hotplugged in while we are in this function,
     * it does not appear in the existing visit_cpus mask
     * which came from the sched_group pointer of the
     * sched_domain pointed at by sd_ea for either the prev
     * or next cpu and was dereferenced in __energy_diff.
     * Since we will dereference sd_scs later as we iterate
     * through the CPUs we expect to visit, new CPUs can
     * be present which are not in the visit_cpus mask.
     * Guard this with cpu_count.
     */
    cpu_count = cpumask_weight(&visit_cpus);
    /*根据sg_top顶层sd，找到需要计算的cpu集合visit_cpus，逐个遍历其中每一个
    cpu,这一套复杂的循环算法计算下来，其实就计算了几个power，以cpu0-cpu3为例
    :4个底层sg的power + 1个顶层sg的power*/
    while (!cpumask_empty(&visit_cpus)) {
        struct sched_group *sg_shared_cap = NULL;
        /*选取visit_cpus中的第一个cpu*/
        int cpu = cpumask_first(&visit_cpus);
        struct sched_domain *sd;

        /*
         * Is the group utilization affected by cpus outside this
         * sched_group?
         * This sd may have groups with cpus which were not present
         * when we took visit_cpus.
         */
        sd = rcu_dereference(per_cpu(sd_scs, cpu));

        if (sd && sd->parent)
            sg_shared_cap = sd->parent->groups;
        /*从底层到顶层逐个遍历cpu所在的sd*/
        for_each_domain(cpu, sd) {
            struct sched_group *sg = sd->groups;
            /*如果是顶层sd，只会计算一个sg*/
            /* Has this sched_domain already been visited? */
            if (sd->child && group_first_cpu(sg) != cpu)
                break;
            /*逐个遍历该层次sg链表所在sg*/
            do {
                unsigned long group_util;
                int sg_busy_energy, sg_idle_energy;
                int cap_idx, idle_idx;

```

```

    if (sg_shared_cap && sg_shared_cap->group_weight >=
sg->group_weight)
        eenv->sg_cap = sg_shared_cap;
    else
        eenv->sg_cap = sg;
    /*根据eenv指示的负载变化，找出满足该sg中最大负载cpu的
    capacity_index*/
    cap_idx = find_new_capacity(eenv, sg->sge);

    if (sg->group_weight == 1) {
        /* Remove capacity of src CPU (before task move) */
        if (eenv->trg_cpu == eenv->src_cpu &&
cpumask_test_cpu(eenv->src_cpu,
sched_group_cpus(sg))) {
            eenv->cap.before = sg->sge->cap_states[cap_idx].cap;
            eenv->cap.delta -= eenv->cap.before;
        }
        /* Add capacity of dst CPU (after task move) */
        if (eenv->trg_cpu == eenv->dst_cpu &&
cpumask_test_cpu(eenv->dst_cpu,
sched_group_cpus(sg))) {
            eenv->cap.after = sg->sge->cap_states[cap_idx].cap;
            eenv->cap.delta += eenv->cap.after;
        }
    }
    /*找出sg所有cpu中最小的idle index*/
    idle_idx = group_idle_state(eenv, sg);
    /*累加sg中所有cpu的相对负载，
    最大负载为sg->sge->cap_states[eenv->cap_idx].cap*/
    group_util = group_norm_util(eenv, sg);
    /*计算power = busy_power + idle_power*/
    sg_busy_energy = (group_util *
sg->sge->cap_states[cap_idx].power);
    sg_idle_energy = ((SCHED_LOAD_SCALE-group_util)
        * sg->sge->idle_states[idle_idx].power);

    total_energy += sg_busy_energy + sg_idle_energy;

    if (!sd->child) {
        /*
        * cpu_count here is the number of
        * cpus we expect to visit in this
        * calculation. If we race against
        * hotplug, we can have extra cpus
        * added to the groups we are
        * iterating which do not appear in
        * the visit_cpus mask. In that case
        * we are not able to calculate energy
        * without restarting so we will bail
        * out and use prev_cpu this time.
        */
    }

```



```

•         * cpu is a valid SD_WAKE_AFFINE target.
•         */
•         if (want_affine && (tmp->flags & SD_WAKE_AFFINE) &&
•             cpumask_test_cpu(prev_cpu, sched_domain_span(tmp))) {
•             affine_sd = tmp;
•             break;
•         }
•
•         if (tmp->flags & sd_flag)
•             sd = tmp;
•         else if (!want_affine)
•             break;
•     }
•
•     if (affine_sd) {
•         sd = NULL; /* Prefer wake_affine over balance flags */
•         if (cpu != prev_cpu && wake_affine(affine_sd, p, prev_cpu, sync))
•             new_cpu = cpu;
•     }
•
•     if (sd && !(sd_flag & SD_BALANCE_FORK)) {
•         /*
•          * We're going to need the task's util for capacity_spare_wake
•          * in find_idlest_group. Sync it up to prev_cpu's
•          * last_update_time.
•          */
•         sync_entity_load_avg(&p->se);
•     }
•
•     if (!sd) {
•         if (sd_flag & SD_BALANCE_WAKE) /* XXX always ? */
•             new_cpu = select_idle_sibling(p, prev_cpu, new_cpu);
•     } else {
•         new_cpu = find_idlest_cpu(sd, p, cpu, prev_cpu, sd_flag);
•     }
•     rcu_read_unlock();
•
•     return new_cpu;
• }

```

2 当挑选的cpu与当前唤醒进程所在的cpu不同时,怎么处理?

```

•     cpu = select_task_rq(p, p->wake_cpu, SD_BALANCE_WAKE, wake_flags,
•                           sibling_count_hint);
•     if (task_cpu(p) != cpu) {
•         wake_flags |= WF_MIGRATED;
•         set_task_cpu(p, cpu);
•     }
•     .....
•
• void set_task_cpu(struct task_struct *p, unsigned int new_cpu)
• {

```

```

• #ifdef CONFIG_SCHED_DEBUG
•     /*
•      * We should never call set_task_cpu() on a blocked task,
•      * ttwu() will sort out the placement.
•      */
•     WARN_ON_ONCE(p->state != TASK_RUNNING && p->state != TASK_WAKING &&
•                 !p->on_rq);
•
• #ifdef CONFIG_LOCKDEP
•     /*
•      * The caller should hold either p->pi_lock or rq->lock, when changing
•      * a task's CPU. ->pi_lock for waking tasks, rq->lock for runnable
      tasks.
•      *
•      * sched_move_task() holds both and thus holding either pins the cgroup,
•      * see task_group().
•      *
•      * Furthermore, all task_rq users should acquire both locks, see
•      * task_rq_lock().
•      */
•     WARN_ON_ONCE(debug_locks && !(lockdep_is_held(&p->pi_lock) ||
•                                     lockdep_is_held(&task_rq(p)->lock)));
• #endif
• #endif
•
•     trace_sched_migrate_task(p, new_cpu);
•
•     if (task_cpu(p) != new_cpu) {
•         if (p->sched_class->migrate_task_rq)
•             p->sched_class->migrate_task_rq(p);
•         p->se.nr_migrations++;
•         perf_event_task_migrate(p);
•
•         walt_fixup_busy_time(p, new_cpu);
•     }
•
•     __set_task_cpu(p, new_cpu);
• }
•
• /*
•  * Called immediately before a task is migrated to a new cpu; task_cpu(p)
  and
•  * cfs_rq_of(p) references at time of call are still valid and identify the
•  * previous cpu. However, the caller only guarantees p->pi_lock is held; no
•  * other assumptions, including the state of rq->lock, should be made.
•  */
• static void migrate_task_rq_fair(struct task_struct *p)
• {
•     /*
•      * We are supposed to update the task to "current" time, then its up to
      date
•      * and ready to go to new CPU/cfs_rq. But we have difficulty in getting
•      * what current time is, so simply throw away the out-of-date time. This

```

```

•      * will result in the wakee task is less decayed, but giving the wakee
more
•      * load sounds not bad.
•      */
•      /*重新计算在新cpu上的调度实体的负载*/
•      remove_entity_load_avg(&p->se);
•      /*重置新的调度实体的负载的最后更新时间和调度实体的执行时间*/
•      /* Tell new CPU we are migrated */
•      p->se.avg.last_update_time = 0;
•
•      /* We have migrated, no longer consider this task hot */
•      p->se.exec_start = 0;
•  }
•  /*
•   * Synchronize entity load avg of dequeued entity without locking
•   * the previous rq.
•   */
•  void sync_entity_load_avg(struct sched_entity *se)
•  {
•      struct cfs_rq *cfs_rq = cfs_rq_of(se);
•      u64 last_update_time;
•
•      last_update_time = cfs_rq_last_update_time(cfs_rq);
•      __update_load_avg(last_update_time, cpu_of(rq_of(cfs_rq)), &se->avg, 0,
0, NULL);
•  }
•
•  /*
•   * Task first catches up with cfs_rq, and then subtract
•   * itself from the cfs_rq (task must be off the queue now).
•   */
•  void remove_entity_load_avg(struct sched_entity *se)
•  {
•      struct cfs_rq *cfs_rq = cfs_rq_of(se);
•
•      /*
•       * tasks cannot exit without having gone through wake_up_new_task() ->
•       * post_init_entity_util_avg() which will have added things to the
•       * cfs_rq, so we can remove unconditionally.
•       *
•       * Similarly for groups, they will have passed through
•       * post_init_entity_util_avg() before unregister_sched_fair_group()
•       * calls this.
•       */
•
•      sync_entity_load_avg(se);
•      atomic_long_add(se->avg.load_avg, &cfs_rq->removed_load_avg);
•      atomic_long_add(se->avg.util_avg, &cfs_rq->removed_util_avg);
•  }
•

```

3 How task p enqueue?

入队操作:ttwu_queue(p, cpu);

```

•  static void ttwu_queue(struct task_struct *p, int cpu)

```

```

• {
•     struct rq *rq = cpu_rq(cpu);
•
•     #if defined(CONFIG_SMP)
•         if (sched_feat(TTWU_QUEUE) && !cpus_share_cache(smp_processor_id(),
•             cpu)) {
•             sched_clock_cpu(cpu); /* sync clocks x-cpu */
•             ttwu_queue_remote(p, cpu);
•             return;
•         }
•     #endif
•
•     raw_spin_lock(&rq->lock);
•     lockdep_pin_lock(&rq->lock);
•     ttwu_do_activate(rq, p, 0);
•     lockdep_unpin_lock(&rq->lock);
•     raw_spin_unlock(&rq->lock);
• }
•
• static void
• ttwu_do_activate(struct rq *rq, struct task_struct *p, int wake_flags)
• {
•     lockdep_assert_held(&rq->lock);
•
•     #ifdef CONFIG_SMP
•         if (p->sched_contributes_to_load)
•             rq->nr_uninterruptible--;
•     #endif
•
•     ttwu_activate(rq, p, ENQUEUE_WAKEUP | ENQUEUE_WAKING);
•     ttwu_do_wakeup(rq, p, wake_flags);
• }
•
• static inline void ttwu_activate(struct rq *rq, struct task_struct *p, int
• en_flags)
• {
•     /* 实际入队的核心函数,同时更新task的vruntime并插入rb tree*/
•     activate_task(rq, p, en_flags);
•     p->on_rq = TASK_ON_RQ_QUEUED;
•
•     /* if a worker is waking up, notify workqueue */
•     if (p->flags & PF_WQ_WORKER)
•         wq_worker_waking_up(p, cpu_of(rq));
• }
•
• /*
•  * Mark the task runnable and perform wakeup-preemption.
•  */
• static void
• ttwu_do_wakeup(struct rq *rq, struct task_struct *p, int wake_flags)
• {
•     check_preempt_curr(rq, p, wake_flags);
•     /* 修改task的状态为running状态,即task正在cpu上运行*/
•     p->state = TASK_RUNNING;
•     trace_sched_wakeup(p);

```

```

•
• #ifdef CONFIG_SMP
•     /*cfs没有定义*/
•     if (p->sched_class->task_woken) {
•         /*
•          * Our task @p is fully woken up and running; so its safe to
•          * drop the rq->lock, hereafter rq is only used for statistics.
•          */
•         lockdep_unpin_lock(&rq->lock);
•         p->sched_class->task_woken(rq, p);
•         lockdep_pin_lock(&rq->lock);
•     }
•     /*如果之前rq处于idle状态,即cpu处于idle状态,则修正rq的idle时间戳和判决idle时间*/
•     if (rq->idle_stamp) {
•         u64 delta = rq_clock(rq) - rq->idle_stamp;
•         u64 max = 2*rq->max_idle_balance_cost;//max=1ms
•
•         update_avg(&rq->avg_idle, delta);
•
•         if (rq->avg_idle > max)
•             rq->avg_idle = max;
•
•         rq->idle_stamp = 0;
•     }
• #endif
• }

```

3.1 核心函数activate_task分析

```

• void activate_task(struct rq *rq, struct task_struct *p, int flags)
• {
•     /*如果进程被置为TASK_UNINTERRUPTIBLE状态的话,减少处于uninterruptible状态的进程
•     数量,因为当前是进程处于唤醒运行阶段*/
•     if (task_contributes_to_load(p))
•         rq->nr_uninterruptible--;
•     /*入队操作*/
•     enqueue_task(rq, p, flags);
• }
•
• #define task_contributes_to_load(task) \
•     ((task->state & TASK_UNINTERRUPTIBLE) != 0 && \
•      (task->flags & PF_FROZEN) == 0 && \
•      (task->state & TASK_NOLOAD) == 0)
•
• static inline void enqueue_task(struct rq *rq, struct task_struct *p, int
flags)
• {
•     update_rq_clock(rq);
•     /*flag=0x3 & 0x10 = 0*/
•     if (!(flags & ENQUEUE_RESTORE))
•         /*更新task p入队的时间戳*/
•         sched_info_queued(rq, p);
• #ifdef CONFIG_INTEL_DWS //没有定义
•     if (sched_feat(INTEL_DWS))
•         update_rq_runnable_task_avg(rq);
• #endif
• }

```

```

• /*callback enqueue_task_fair函数*/
• p->sched_class->enqueue_task(rq, p, flags);
• }

```

核心函数enqueue_task_fair解析如下:

```

• /*
•  * The enqueue_task method is called before nr_running is
•  * increased. Here we update the fair scheduling stats and
•  * then put the task into the rbtree:
•  */
• static void
• enqueue_task_fair(struct rq *rq, struct task_struct *p, int flags)
• {
•     struct cfs_rq *cfs_rq;
•     struct sched_entity *se = &p->se;
• #ifdef CONFIG_SMP
•     /*task_new = 0x3 & 0x20 = 0*/
•     int task_new = flags & ENQUEUE_WAKEUP_NEW;
• #endif
•     /*进程p开始运行,并将进程p的运行时间累加到整个rq的cumulative_runnable_avg时间
•     上,作为rq的负载值*/
•     walt_inc_cumulative_runnable_avg(rq, p);
•
•     /*
•     * Update SchedTune accounting.
•     *
•     * We do it before updating the CPU capacity to ensure the
•     * boost value of the current task is accounted for in the
•     * selection of the OPP.
•     *
•     * We do it also in the case where we enqueue a throttled task;
•     * we could argue that a throttled task should not boost a CPU,
•     * however:
•     * a) properly implementing CPU boosting considering throttled
•     *    tasks will increase a lot the complexity of the solution
•     * b) it's not easy to quantify the benefits introduced by
•     *    such a more complex solution.
•     * Thus, for the time being we go for the simple solution and boost
•     * also for throttled RQs.
•     */
•     /*主要根据这个task的属性是否需要update task group的boost参数*/
•     schedtune_enqueue_task(p, cpu_of(rq));
•
•     /*
•     * If in_iowait is set, the code below may not trigger any cpufreq
•     * utilization updates, so do it here explicitly with the IOWAIT flag
•     * passed.
•     */
•     /*如果进程p是一个iowait的进程,则进行cpu频率调整*/
•     if (p->in_iowait)
•         cpufreq_update_util(rq, SCHED_CPUFREQ_IOWAIT);
•
•     for_each_sched_entity(se) {
•         if (se->on_rq)
•             break;
•

```

```

•      cfs_rq = cfs_rq_of(se);
•      walt_inc_cfs_cumulative_runnable_avg(cfs_rq, p);
•      enqueue_entity(cfs_rq, se, flags);
•
•      /*
•       * end evaluation on encountering a throttled cfs_rq
•       *
•       * note: in the case of encountering a throttled cfs_rq we will
•       * post the final h_nr_running increment below.
•       */
•      if (cfs_rq_throttled(cfs_rq))
•          break;
•      cfs_rq->h_nr_running++;
•
•      flags = ENQUEUE_WAKEUP;
•  }
•
•  for_each_sched_entity(se) {
•      cfs_rq = cfs_rq_of(se);
•      cfs_rq->h_nr_running++;
•      walt_inc_cfs_cumulative_runnable_avg(cfs_rq, p);
•
•      if (cfs_rq_throttled(cfs_rq))
•          break;
•
•      update_load_avg(se, UPDATE_TG);
•      update_cfs_shares(se);
•  }
•
•  if (!se)
•      add_nr_running(rq, 1);
•
•  #ifdef CONFIG_SMP
•      if (!se) {
•          struct sched_domain *sd;
•
•          rcu_read_lock();
•          sd = rcu_dereference(rq->sd);
•          if (!task_new && sd) {
•              if (cpu_overutilized(rq->cpu))
•                  set_sd_overutilized(sd);
•              if (rq->misfit_task && sd->parent)
•                  set_sd_overutilized(sd->parent);
•          }
•          rcu_read_unlock();
•      }
•
•  #endif /* CONFIG_SMP */
•      hrtick_update(rq);
•  }

```

后面的流程与新创建进程的流程一致了：

<https://blog.csdn.net/wukongmingjing/article/details/82466628>

