


scheduler里面这个负载的概念可能被误解为cpu占用率，但是在调度里面这个有比较大的偏差。scheduler不使用cpu占用率来评估负载，而是使用平均时间runnable的数量来评估负载。Scheduler也分了几个层级来计算负载：

- entity级负载计算：update_load_avg()
- cpu级负载计算：update_cpu_load_active()
- 系统级负载计算：calc_global_load_tick()

第一个是调度实体的load,即sched_entity的load,根据PELT算法实现的,算法逻辑如下:



PELT(Per Entity Load Tracing)算法概述

从上面示意图可以看到,task runtime是 $\delta = \delta_1 + \delta_2 + \delta_3$ 之和

- δ 数值依赖真实task的运行时间,是总的运行时间
- last update time是task load是上次更新的最后时间(第一个红色箭头)
- now是task load更新的当前时间(第二个红色箭头)
- 1ms表示1024us的颗粒度,由于kernel对于除法效率较低不能使用小数,所以1ms直接转化为1024us,好做乘法和位移运算,真的很巧妙

示意图的目的就是追踪三个时间段(phase1/phase2/phase3)的load,来计算now时刻的load,周而复始.

PELT算法

Phase1阶段怎么计算load

1. 计算 δ_1 的period:

$$\delta_1 = 1024 - \text{Period_contrib1} \quad (< 1024\text{us})$$
2. load_sum被刻度化通过当前cpu频率和se的权重:

$$\delta_1 = \delta_1 * \text{scale_freq}$$

$$\text{load_sum} += \text{weight} * \delta_1$$
3. load_util被cpu的capacity刻度化

$$\text{util_sum} += \text{scale_cpu} * \delta_1;$$

Phase2阶段怎么计算load的:

1. 计算 δ_2 的period

$$\text{periods} = \delta_2 / 1024 \text{ (即存在有多少个1ms)}$$
2. 衰减phase1的load

$$\text{load_sum} += \text{decay_load}(\text{load_sum}, \text{periods} + 1)$$

$$\text{util_sum} += \text{decay_load}(\text{util_sum}, \text{periods} + 1)$$

3. 衰减阶段phase2的load

```
load_sum += __contrib(periods) * scale_freq  
util_sum += __contrib(periods) * scale_freq * scale_cpu
```

Phase3计算怎么计算load:

1. 计算剩余周期(<1ms,<1024us)
period_contrib2 = delta3 % 1024
2. load_sum被当前权重和频率刻度化
load_sum += weight * scale_freq * period_contrib2
3. util_sum被当前频率和当前cpu capacity刻度化
util_sum += scale_cpu * scale_freq * period_contrib2

上面是这个算法的精髓以及思路,下面讲解decay_load和__contrib怎么计算的

decay_load:

对于每一个period(大小为LOAD_AVG_PERIOD=32ms),这个load将衰减0.5,因此根据当前period,load被衰减方式如下:

1. $load = (load \gg (n/period)) * y^{(n \% period)}$
2. 并且 $y^{(n \% period)}$ 可以看成runnable_avg_yN_inv[n]的数值

在kernel中查表即可:

```
/* Precomputed fixed inverse multiplies for multiplication by y^n */  
static const u32 runnable_avg_yN_inv[] = {  
    0xffffffff, 0xfa83b2da, 0xf5257d14, 0xfe4b99a, 0xeac0c6e6, 0xe5b906e6,  
    0xe0ccdeeb, 0xdbfb796, 0xd744fcc9, 0xd2a81d91, 0xce248c14, 0xc9b9bd85,  
    0xc5672a10, 0xc12c4cc9, 0xbd08a39e, 0xb8fbaf46, 0xb504f333, 0xb123f581,  
    0xad583ee9, 0xa9a15ab4, 0xa5fed6a9, 0xa2704302, 0x9ef5325f, 0x9b8d39b9,  
    0x9837f050, 0x94f4efa8, 0x91c3d373, 0x8ea4398a, 0x8b95c1e3, 0x88980e80,  
    0x85aac367, 0x82cd8698,  
};  
  
#define LOAD_AVG_PERIOD 32  
#define LOAD_AVG_MAX 47742 /* maximum possible load avg */  
#define LOAD_AVG_MAX_N 345 /* number of full periods to produce LOAD_AVG_MAX */
```

其实现代码如下:

```
/*  
 * Approximate:  
 *   val * y^n,   where y^32 ~= 0.5 (~1 scheduling period)  
 */  
static __always_inline u64 decay_load(u64 val, u64 n)  
{  
    unsigned int local_n;  
  
    if (!n)  
        return val;  
    else if (unlikely(n > LOAD_AVG_PERIOD * 63))  
        return 0;  
  
    /* after bounds checking we can collapse to 32-bit */  
    local_n = n;  
  
    /*
```

```

•   * As  $y^{\text{PERIOD}} = 1/2$ , we can combine
•   *    $y^n = 1/2^{(n/\text{PERIOD})} * y^{(n\% \text{PERIOD})}$ 
•   * With a look-up table which covers  $y^n$  ( $n < \text{PERIOD}$ )
•   *
•   * To achieve constant time decay_load.
•   */ /*LOAD_AVG_PERIOD = 32*/
•   if (unlikely(local_n >= LOAD_AVG_PERIOD)) {
•       val >>= local_n / LOAD_AVG_PERIOD;
•       local_n %= LOAD_AVG_PERIOD;
•   }
•   /*正好符合:load = (load >> (n/period)) * y^(n%period)计算方式*/
•   val = mul_u64_u32_shr(val, runnable_avg_yN_inv[local_n], 32);
•   return val;
•   }

```

__contrib:

1. if $\text{period} \leq \text{LOAD_AVG_PERIOD}(32\text{ms}, 32 * 1024\text{us})$
 $\text{load} = 1024 + 1024*y + 1024*y^2 + \dots + 1024*y^{\text{period}}$
2. if $\text{period} > \text{LOAD_AVG_MAX_N}(345\text{ms})$
 $\text{load} = \text{LOAD_AVG_MAX}(47742)$
3. if $\text{period} \in (32, 345]$, 即每个 LOAD_AVG_PERIOD 周期衰减累加
do
 $\text{load} /= 2$
 $\text{load} += 1024 + 1024*y + 1024*y^2 + \dots + 1024*y^{\text{LOAD_AVG_PERIOD}}$
 $n -= \text{period}$
while($n > \text{LOAD_AVG_PERIOD}$)
4. $1024 + 1024*y + 1024*y^2 + \dots + 1024*y^{32} = \text{runnable_avg_yN_sum}[32]$
decay_load()只是计算 y^n , 而__contrib是计算一个对比队列的和: $y + y^2 + y^3 \dots + y^n$. 计算方式如下:
runnable_avg_yN_sum[]数组是使用查表法来计算 $n=32$ 位内的等比队列求和:
 $\text{runnable_avg_yN_sum}[1] = y^1 * 1024 = 0.978520621 * 1024 = 1002$
 $\text{runnable_avg_yN_sum}[2] = (y^1 + y^2) * 1024 = 1982$
...
 $\text{runnable_avg_yN_sum}[32] = (y^1 + y^2 \dots + y^{32}) * 1024 = 23371$

实现代码和查表数据如下:

```

•   static u32 __compute_runnable_contrib(u64 n)
•   {
•       u32 contrib = 0;
•
•       if (likely(n <= LOAD_AVG_PERIOD))
•           return runnable_avg_yN_sum[n];
•       else if (unlikely(n >= LOAD_AVG_MAX_N))
•           return LOAD_AVG_MAX;
•
•       /* Compute \Sum k^n combining precomputed values for k^i, \Sum k^j */
•       do {
•           contrib /= 2; /* y^LOAD_AVG_PERIOD = 1/2 */
•           contrib += runnable_avg_yN_sum[LOAD_AVG_PERIOD];
•
•           n -= LOAD_AVG_PERIOD;
•       }
•   }

```

```

•     } while (n > LOAD_AVG_PERIOD);
•
•     contrib = decay_load(contrib, n);
•     return contrib + runnable_avg_yN_sum[n];
• }
• /*
•  * Precomputed \Sum y^k { 1<=k<=n }. These are floor(true_value) to prevent
•  * over-estimates when re-combining.
•  */
• static const u32 runnable_avg_yN_sum[] = {
•     0, 1002, 1982, 2941, 3880, 4798, 5697, 6576, 7437, 8279, 9103,
•     9909, 10698, 11470, 12226, 12966, 13690, 14398, 15091, 15769, 16433, 17082,
•     17718, 18340, 18949, 19545, 20128, 20698, 21256, 21802, 22336, 22859, 23371,
• };

```

针对__contrib第二点当period>345的时候,load变成一个常数怎么理解的?

即**load = LOAD_AVG_MAX (47742)**,我们简单来证明以下:

设一个等比数列的首项是a1,公比是y, 数列前n项和是Sn, 当公比不为1时

$$S_n = a_1 + a_1y + a_1y^2 + \dots + a_1y^{(n-1)}$$

将这个式子两边同时乘以公比y, 得

$$yS_n = a_1y + a_1y^2 + \dots + a_1y^{(n-1)} + a_1y^n$$

两式相减, 得

$$(1-y) S_n = a_1 - a_1y^n$$

所以, 当公比不为1时, 等比数列的求和公式:

$$S_n = a_1(1-y^n)/(1-y)$$

对于一个无穷递降数列, 数列的公比小于1, 当上式得n趋向于正无穷大时, 分子括号中的值趋近于1, 取极限即得无穷递减数列求和公式:


$$S = a_1/(1-y)$$

由于要使 $y^{32} = 0.5$, 那么经过计算之后, $y \approx 0.9785$ ($0.9785^{32} \approx 0.498823$)

所以对于period > **LOAD_AVG_MAX_N(345)**, 等比数列求和数值如下:

$$sn = a_1(1-y^n)/(1-y) = 1024 * (1 - 0.9785^n) / (1 - 0.9785) \approx 1002 * (1 - 0.9785^n)$$

画出曲线图如下:



从当n趋于一个数值, 当n增大, 等比数列之后增加几乎可以忽略, 并且无穷大 ∞ , 则等比数列之和为 $a_1/(1-y) = 1024/(1-0.9785) \approx 47627.9069988967$, 与47742数值差不多. 上面说明了原理, 下面就是实际的代码分析了.

那么上面的两个表格runnable_avg_yN_inv和runnable_avg_yN_sum是怎么计算的, 下面是一个通过C语言计算的小程序:

```
#include <stdio.h>
#include <math.h>
#if 1
#define N 32
#define WMULT_SHIFT 32

const long WMULT_CONST = ((1UL << N) - 1);
double y;

long runnable_avg_yN_inv[N];
void calc_mult_inv()
{
    int i;
    double yn = 0;

    printf("inverses\n");
    for (i = 0; i < N; i++) {
        yn = (double)WMULT_CONST * pow(y, i);
        runnable_avg_yN_inv[i] = yn;
        printf("%2d: 0x%8lx\n", i, runnable_avg_yN_inv[i]);
    }
    printf("\n");
}
```

```

•
• long mult_inv(long c, int n)
• {
•     return (c * runnable_avg_yN_inv[n]) >> WMULT_SHIFT;
• }
•
• void calc_yn_sum(int n)
• {
•     int i;
•     double sum = 0, sum_fl = 0, diff = 0;
•
•     /*
•      * We take the floored sum to ensure the sum of partial sums is never
•      * larger than the actual sum.
•      */
•     printf("sum y^n\n");
•     printf("    %8s %8s %8s\n", "exact", "floor", "error");
•     for (i = 1; i <= n; i++) {
•         sum = (y * sum + y * 1024);
•         sum_fl = floor(y * sum_fl + y * 1024);
•         printf("%2d: %8.0f %8.0f %8.0f\n", i, sum, sum_fl,
•             sum_fl - sum);
•     }
•     printf("\n");
• }
•
• void calc_conv(long n)
• {
•     long old_n;
•     int i = -1;
•
•     printf("convergence (LOAD_AVG_MAX, LOAD_AVG_MAX_N)\n");
•     do {
•         old_n = n;
•         n = mult_inv(n, 1) + 1024;
•         i++;
•     } while (n != old_n);
•     printf("%d> %ld\n", i - 1, n);
•     printf("\n");
• }
•
• #endif
• int main(void)
• {
•     #if 1
•         y = pow(0.5, 1/(double)N);
•         calc_mult_inv();
•         calc_conv(1024);
•         calc_yn_sum(N);
•     #endif
•
•     return 0;
• }

```

PELT Entity级的负载计算

- Entity级的负载计算也称作PELT(Per-Entity Load Tracking)。
- 注意负载计算时使用的时间都是实际运行时间而不是虚拟运行时间vruntime。

过程如下:

```
• scheduler_tick() -> task_tick_fair() -> entity_tick() -> update_load_avg()
• /* Update task and its cfs_rq load average */
• static inline void update_load_avg(struct sched_entity *se, int flags)
• {
•     struct cfs_rq *cfs_rq = cfs_rq_of(se);
•     u64 now = cfs_rq_clock_task(cfs_rq);
•     int cpu = cpu_of(rq_of(cfs_rq));
•     int decayed;
•     void *ptr = NULL;
•
•     /*
•      * Track task load average for carrying it to new CPU after migrated,
and
•      * track group sched_entity load average for task_h_load calc in
migration
•      **cfs load tracing时间已经update,也就是已经初始化过了
•      SKIP_AGE_LOAD是忽略load tracing的flag*/
•     if (se->avg.last_update_time && !(flags & SKIP_AGE_LOAD)) {
•         /*核心函数,即PELT的实现,注意se->on_rq的数值,如果一直在运行的进程,则
•         se->on_rq,load=老负载衰减+新负载,如果是休眠唤醒进程se->on_rq=0,则他们在
•         休眠期间的load不会累加,只有老负载被衰减,睡眠时间不会统计在内,直到task在rq里面*/
•         __update_load_avg(now, cpu, &se->avg,
•             se->on_rq * scale_load_down(se->load.weight),
•             cfs_rq->curr == se, NULL);
•     }
•
•     decayed = update_cfs_rq_load_avg(now, cfs_rq, true);
•     decayed |= propagate_entity_load_avg(se);
•
•     if (decayed && (flags & UPDATE_TG))
•         update_tg_load_avg(cfs_rq, 0);
•
•     if (entity_is_task(se)) {
• #ifdef CONFIG_SCHED_WALT
•         ptr = (void *)&(task_of(se)->ravg);
• #endif
•         trace_sched_load_avg_task(task_of(se), &se->avg, ptr);
•     }
• }
```

核心函数1 __update_load_avg()的实现

我们先明白下面几个参数的含义:

- load_sum
- util_sum

3. load_avg

4. util_avg

上面几个涉及到cfs_rq结构体的成员变量:

```
• struct cfs_rq {
• .....
•     /*
•     * CFS load tracking
•     */
•     struct sched_avg avg;
•     u64 runnable_load_sum;
•     unsigned long runnable_load_avg;
•     .....
• }
• /*
• * The load_avg/util_avg accumulates an infinite geometric series.
• * 1) load_avg factors frequency scaling into the amount of time that a
• * sched_entity is runnable on a rq into its weight. For cfs_rq, it is the
• * aggregated such weights of all runnable and blocked sched_entities.
• * 2) util_avg factors frequency and cpu scaling into the amount of time
• * that a sched_entity is running on a CPU, in the range
• [0..SCHED_LOAD_SCALE].
• * For cfs_rq, it is the aggregated such times of all runnable and
• * blocked sched_entities.
• * The 64 bit load_sum can:
• * 1) for cfs_rq, afford 4353082796 (=2^64/47742/88761) entities with
• * the highest weight (=88761) always runnable, we should not overflow
• * 2) for entity, support any load.weight always runnable
• */
• struct sched_avg {
•     u64 last_update_time, load_sum;
•     u32 util_sum, period_contrib;
•     unsigned long load_avg, util_avg;
• };
```

而且如果知道了load_sum,util_sum,runnable_load_sum,这几个数值除以LOAD_AVG_MAX(**47742**)则就可以直接计算load_avg,util_avg,runnable_load_avg,即:
util_avg = util_sum / **LOAD_AVG_MAX(47742)**.

5. scale_freq:<https://blog.csdn.net/wukongmingjing/article/details/81635383>

6. scale_cpu:<https://blog.csdn.net/wukongmingjing/article/details/81635383>

关键函数的代码如下:

```
• /*
• * We can represent the historical contribution to runnable average as the
• * coefficients of a geometric series. To do this we sub-divide our
• runnable
• * history into segments of approximately 1ms (1024us); label the segment
• that
• * occurred N-ms ago p_N, with p_0 corresponding to the current period, e.g.
• *
• * [<- 1024us ->|<- 1024us ->|<- 1024us ->| ...
• *      p0              p1              p2
• *      (now)          (~1ms ago)    (~2ms ago)
• *
• * Let u_i denote the fraction of p_i that the entity was runnable.
• *
```



```

• * We then designate the fractions u_i as our co-efficients, yielding the
• * following representation of historical load:
• *   u_0 + u_1*y + u_2*y^2 + u_3*y^3 + ...
• *
• * We choose y based on the width of a reasonably scheduling period, fixing:
• *   y^32 = 0.5
• *
• * This means that the contribution to load ~32ms ago (u_32) will be
weighted
• * approximately half as much as the contribution to load within the last ms
• * (u_0).
• *
• * When a period "rolls over" and we have new u_0`, multiplying the previous
• * sum again by y is sufficient to update:
• *   load_avg = u_0` + y*(u_0 + u_1*y + u_2*y^2 + ... )
• *               = u_0 + u_1*y + u_2*y^2 + ... [re-labeling u_i --> u_{i+1}]
• */
• static __always_inline int
• __update_load_avg(u64 now, int cpu, struct sched_avg *sa,
•                  unsigned long weight, int running, struct cfs_rq *cfs_rq)
• {
•     u64 delta, scaled_delta, periods;
•     u32 contrib;
•     unsigned int delta_w, scaled_delta_w, decayed = 0;
•     unsigned long scale_freq, scale_cpu;
•
• #ifdef CONFIG_64BIT_ONLY_CPU
•     struct sched_entity *se;
•     unsigned long load_avg_before = sa->load_avg;
• #endif
•     /*就是示意图中的delta1+delta2+delta3*/
•     delta = now - sa->last_update_time;
•     /*
•      * This should only happen when time goes backwards, which it
•      * unfortunately does during sched clock init when we swap over to TSC.
•      */
•     if ((s64)delta < 0) {
•         sa->last_update_time = now;
•         return 0;
•     }
•
•     /*
•      * Use 1024ns as the unit of measurement since it's a reasonable
•      * approximation of 1us and fast to compute.
•      */
•     delta >= 10;
•     if (!delta)
•         return 0;
•     sa->last_update_time = now;
•     /*scale_freq = (curr_freq << 10)/policy->max*/
•     scale_freq = arch_scale_freq_capacity(NULL, cpu);
•     /*scale_cpu = capacity[cpu],dts获取的,不同cluster capacity不同*/
•     scale_cpu = arch_scale_cpu_capacity(NULL, cpu);
•     trace_sched_contrib_scale_f(cpu, scale_freq, scale_cpu);

```

```

/* delta_w is the amount already accumulated against our next period */
delta_w = sa->period_contrib;
/*表示delta1+delta2大于一个最小刻度1024,如果小于,则就只剩下delta3计算,delta1,
delta2不存在*/
if (delta + delta_w >= 1024) {
    decayed = 1;

    /* how much left for next period will start over, we don't know yet
*/
    sa->period_contrib = 0;

    /*
     * Now that we know we're crossing a period boundary, figure
     * out how much from delta we need to complete the current
     * period and accrue it.
     */
    /*开始Phase1阶段的load_sum 和util_sum的计算*/
    delta_w = 1024 - delta_w;
    scaled_delta_w = cap_scale(delta_w, scale_freq);
    if (weight) {
        sa->load_sum += weight * scaled_delta_w;
        if (cfs_rq) {
            cfs_rq->runnable_load_sum +=
                weight * scaled_delta_w;
        }
    }
    if (running)
        sa->util_sum += scaled_delta_w * scale_cpu;
    /*结束Phase1阶段的load_sum 和util_sum的计算*/
    delta -= delta_w;

    /* Figure out how many additional periods this update spans */
    /*开始Phase2阶段的load_sum 和util_sum的计算,计算阶段Phase2存在多少个1024
    的倍数和余数*/
    periods = delta / 1024;
    delta %= 1024;
    /*对阶段Phase1的load_sum进行衰减*/
    sa->load_sum = decay_load(sa->load_sum, periods + 1);
    if (cfs_rq) {
        /*对阶段Phase1的runnable_load_sum进行衰减*/
        cfs_rq->runnable_load_sum =
            decay_load(cfs_rq->runnable_load_sum, periods + 1);
    }
    /*对Phase1阶段util_sum进行衰减*/
    sa->util_sum = decay_load((u64)(sa->util_sum), periods + 1);
    /*至此,上面已经得到了阶段Phase2衰减前的load_sum,util_sum,
    runnable_load_sum的数值*/
    /* Efficiently calculate \sum (1..n_period) 1024*y^i */

    /*对Phase2的load/util数据进行衰减*/
    contrib = __compute_runnable_contrib(periods);
    contrib = cap_scale(contrib, scale_freq);
    if (weight) {

```

```

    sa->load_sum += weight * contrib;
    if (cfs_rq)
        cfs_rq->runnable_load_sum += weight * contrib;
}
if (running)
    sa->util_sum += contrib * scale_cpu;
}
/*结束Phase2阶段的load_sum 和util_sum的计算*/
/* Remainder of delta accrued against u_0` */
/*开始阶段Phase3的的load/util的计算*/
scaled_delta = cap_scale(delta, scale_freq);
if (weight) {
    sa->load_sum += weight * scaled_delta;
    if (cfs_rq)
        cfs_rq->runnable_load_sum += weight * scaled_delta;
}
if (running)
    sa->util_sum += scaled_delta * scale_cpu;
/*结束阶段Phase3的的load/util的计算*/
/*sa->period_contrib ∈ [0,1024)*/
sa->period_contrib += delta;
/*如果衰减了,则计算load的avg的数值,否则由于颗粒度太小,没有计算的必要*/
if (decayed) {
    sa->load_avg = div_u64(sa->load_sum, LOAD_AVG_MAX);
    if (cfs_rq) {
        cfs_rq->runnable_load_avg =
            div_u64(cfs_rq->runnable_load_sum, LOAD_AVG_MAX);
    }
    sa->util_avg = sa->util_sum / LOAD_AVG_MAX;
}

#ifdef CONFIG_64BIT_ONLY_CPU
    if (!cfs_rq) {
        if (is_sched_avg_32bit(sa)) {
            se = container_of(sa, struct sched_entity, avg);
            cfs_rq_of(se)->runnable_load_avg_32bit +=
                sa->load_avg - load_avg_before;
        }
    }
#endif

    return decayed;
}
/*
 * Approximate:
 *   val * y^n,   where y^32 ≈ 0.5 (~1 scheduling period)
 */
static __always_inline u64 decay_load(u64 val, u64 n)
{
    unsigned int local_n;

    if (!n)
        return val;
    else if (unlikely(n > LOAD_AVG_PERIOD * 63))

```

```

    return 0;

    /* after bounds checking we can collapse to 32-bit */
    local_n = n;
    /*计算公式为:load = (load >> (n/period)) * y^(n%period),如果n是32的整数倍
    , 因为 $2^{32} = 1/2$ , 相当于右移一位计算n有多少个32, 每个32右移一位*/
    /*
    * As  $y^{\text{PERIOD}} = 1/2$ , we can combine
    *  $y^n = 1/2^{(n/\text{PERIOD})} * y^{(n\%\text{PERIOD})}$ 
    * With a look-up table which covers  $y^n$  ( $n < \text{PERIOD}$ )
    *
    * To achieve constant time decay_load.
    */
    if (unlikely(local_n >= LOAD_AVG_PERIOD)) {
        val >>= local_n / LOAD_AVG_PERIOD;
        local_n %= LOAD_AVG_PERIOD;
    }
    /*将val*y^32,转化为val*runnable_avg_yN_inv[n%LOAD_AVG_PERIOD]>>32*/
    val = mul_u64_u32_shr(val, runnable_avg_yN_inv[local_n], 32);
    return val;
}

/*
 * For updates fully spanning n periods, the contribution to runnable
 * average will be: \Sum 1024*y^n
 *
 * We can compute this reasonably efficiently by combining:
 *  $y^{\text{PERIOD}} = 1/2$  with precomputed \Sum 1024*y^n {for n < PERIOD}
 */
static u32 __compute_runnable_contrib(u64 n)
{
    u32 contrib = 0;

    if (likely(n <= LOAD_AVG_PERIOD))
        return runnable_avg_yN_sum[n];
    else if (unlikely(n >= LOAD_AVG_MAX_N))
        return LOAD_AVG_MAX;
    /*如果n>32, 计算32的整数部分*/
    /* Compute \Sum k^n combining precomputed values for k^i, \Sum k^j */
    do {
        /*每整数32的衰减就是0.5*/
        contrib /= 2; /*  $y^{\text{LOAD\_AVG\_PERIOD}} = 1/2$  */
        contrib += runnable_avg_yN_sum[LOAD_AVG_PERIOD];

        n -= LOAD_AVG_PERIOD;
    } while (n > LOAD_AVG_PERIOD);

    /*将整数部分对余数n进行衰减*/
    contrib = decay_load(contrib, n);
    /*剩余余数n, 使用查表法计算*/
    return contrib + runnable_avg_yN_sum[n];
}

```

核心函数2 update_cfs_rq_load_avg()的实现

```
• /**
•  * update_cfs_rq_load_avg - update the cfs_rq's load/util averages
•  * @now: current time, as per cfs_rq_clock_task()
•  * @cfs_rq: cfs_rq to update
•  * @update_freq: should we call cfs_rq_util_change() or will the call do so
•  *
•  * The cfs_rq avg is the direct sum of all its entities (blocked and
runnable)
•  * avg. The immediate corollary is that all (fair) tasks must be attached,
see
•  * post_init_entity_util_avg().
•  *
•  * cfs_rq->avg is used for task_h_load() and update_cfs_share() for example.
•  *
•  * Returns true if the load decayed or we removed load.
•  *
•  * Since both these conditions indicate a changed cfs_rq->avg.load we should
•  * call update_tg_load_avg() when this function returns true.
•  */
• static inline int
• update_cfs_rq_load_avg(u64 now, struct cfs_rq *cfs_rq, bool update_freq)
• {
•     struct sched_avg *sa = &cfs_rq->avg;
•     int decayed, removed = 0, removed_util = 0;
•     /*是否设置了remove_load_avg和remove_util_avg,如果设置了就修正之前计算的
load/util数值*/
•     if (atomic_long_read(&cfs_rq->removed_load_avg)) {
•         s64 r = atomic_long_xchg(&cfs_rq->removed_load_avg, 0);
•         sub_positive(&sa->load_avg, r);
•         sub_positive(&sa->load_sum, r * LOAD_AVG_MAX);
•         removed = 1;
•         set_tg_cfs_propagate(cfs_rq);
•     }
•
•     if (atomic_long_read(&cfs_rq->removed_util_avg)) {
•         long r = atomic_long_xchg(&cfs_rq->removed_util_avg, 0);
•         sub_positive(&sa->util_avg, r);
•         sub_positive(&sa->util_sum, r * LOAD_AVG_MAX);
•         removed_util = 1;
•         set_tg_cfs_propagate(cfs_rq);
•     }
•     /*对校准后的load进行重新计算*/
•     decayed = __update_load_avg(now, cpu_of(rq_of(cfs_rq)), sa,
•         scale_load_down(cfs_rq->load.weight), cfs_rq->curr != NULL, cfs_rq);
•
•     #ifndef CONFIG_64BIT
•         smp_wmb();
•         cfs_rq->load_last_update_time_copy = sa->last_update_time;
•     #endif
•
•     /* Trace CPU load, unless cfs_rq belongs to a non-root task_group */
```

```

•   if (cfs_rq == &rq_of(cfs_rq)->cfs)
•       trace_sched_load_avg_cpu(cpu_of(rq_of(cfs_rq)), cfs_rq);
•   /*如果为true,则调用schedutil governor进行频率的调整!!!*/
•   if (update_freq)
•       cfs_rq_util_change(cfs_rq);
•
•   return decayed || removed;
• }

```

update_load_avg剩下的函数执行如下:

- propagate_entity_load_avg,更新调度实体本身自己的load/util信息.如果是一个进程则不需要propagate处理.
- 根据decayed的数值和需要更新进程组信息,则调用update_tg_load_avg,更新task_group信息

CPU级的负载计算update_cpu_load_active(rq)

__update_load_avg()是计算se/cfs_rq级别的负载, 在cpu级别linux使用update_cpu_load_active(rq)来计算整个cpu->rq负载的变化趋势。计算也是周期性的, 周期为TICK(时间不固定,由于是tickless系统)。

```

•   scheduler_tick()----->
•   /*
•   * Called from scheduler_tick()
•   */
•   void update_cpu_load_active(struct rq *this_rq)
•   { /*获取cfs_rq的runnable_load_avg的数值*/
•       unsigned long load = weighted_cpuload(cpu_of(this_rq));
•       /*
•       * See the mess around update_idle_cpu_load() / update_cpu_load_nohz().
•       */ /*设置更新rq load的时间戳*/
•       this_rq->last_load_update_tick = jiffies;
•       /核心函数*/
•       __update_cpu_load(this_rq, load, 1);
•   }
•   /* Used instead of source_load when we know the type == 0 */
•   static unsigned long weighted_cpuload(const int cpu)
•   {
•       return cfs_rq_runnable_load_avg(&cpu_rq(cpu)->cfs);
•   }
•
•   static inline unsigned long cfs_rq_runnable_load_avg(struct cfs_rq *cfs_rq)
•   {
•       return cfs_rq->runnable_load_avg;
•   }
•
•   /*
•   * Update rq->cpu_load[] statistics. This function is usually called every
•   * scheduler tick (TICK_NSEC). With tickless idle this will not be called
•   * every tick. We fix it up based on jiffies.
•   */
•   static void __update_cpu_load(struct rq *this_rq, unsigned long this_load,

```

```

    unsigned long pending_updates)
{
    int i, scale;
    /*统计数据使用*/
    this_rq->nr_load_updates++;

    /* Update our load: */
    /*将当前最新的load,更新在cpu_load[0]中*/
    this_rq->cpu_load[0] = this_load; /* Fasttrack for idx 0 */
    for (i = 1, scale = 2; i < CPU_LOAD_IDX_MAX; i++, scale += scale) {
        unsigned long old_load, new_load;

        /* scale is effectively 1 << i now, and >> i divides by scale */

        old_load = this_rq->cpu_load[i];
        /*对old_load进行衰减.果因为进入noHZ模式, 有pending_updates个tick没有更新, 先老化原有负载*/
        old_load = decay_load_missed(old_load, pending_updates - 1, i);
        new_load = this_load;
        /*
         * Round up the averaging division if load is increasing. This
         * prevents us from getting stuck on 9 if the load is 10, for
         * example.
         */
        if (new_load > old_load)
            new_load += scale - 1;
        /*cpu_load的计算公式 */
        this_rq->cpu_load[i] = (old_load * (scale - 1) + new_load) >> i;
    }
    /*更新rq的age_stamp时间戳, 即rq从cpu启动到现在存在的时间(包括idle和running时间)
    ,同时更新rq里面rt_avg负载,即每个周期(500ms)衰减一半*/
    sched_avg_update(this_rq);
}

void sched_avg_update(struct rq *rq)
{
    s64 period = sched_avg_period();

    while ((s64) (rq_clock(rq) - rq->age_stamp) > period) {
        /*
         * Inline assembly required to prevent the compiler
         * optimising this loop into a divmod call.
         * See __iter_div_u64_rem() for another example of this.
         */
        asm("" : "+rm" (rq->age_stamp));
        rq->age_stamp += period;
        rq->rt_avg /= 2;
    }
}

```

码注释中详细解释了cpu_load的计算方法：

- 每个tick计算不同idx时间等级的load，计算公式： $load = (2^{idx} - 1) / 2^{idx} * load + 1 / 2^{idx} * cur_load$
- 如果cpu因为noHZ错过了(n-1)个tick的更新，那么计算load要分两步：
 1. 首先老化(decay)原有的load： $load = ((2^{idx} - 1) / 2^{idx})^{(n-1)} * load$
 2. 再按照一般公式计算load： $load = (2^{idx} - 1) / 2^{idx} * load + 1 / 2^{idx} * cur_load$

- 为了decay的加速计算，设计了decay_load_missed()查表法计算：

```

● /*
●  * The exact cpuload at various idx values, calculated at
every tick would be
●  * load = (2^idx - 1) / 2^idx * load + 1 / 2^idx * cur_load
●  *
●  * If a cpu misses updates for n-1 ticks (as it was idle) and
update gets called
●  * on nth tick when cpu may be busy, then we have:
●  * load = ((2^idx - 1) / 2^idx)^(n-1) * load
●  * load = (2^idx - 1) / 2^idx * load + 1 / 2^idx * cur_load
●  *
●  * decay_load_missed() below does efficient calculation of
●  * load = ((2^idx - 1) / 2^idx)^(n-1) * load
●  * avoiding 0..n-1 loop doing load = ((2^idx - 1) / 2^idx) *
load
●  *
●  * The calculation is approximated on a 128 point scale.
●  * degrade_zero_ticks is the number of ticks after which load
at any
●  * particular idx is approximated to be zero.
●  * degrade_factor is a precomputed table, a row for each load
idx.
●  * Each column corresponds to degradation factor for a power
of two ticks,
●  * based on 128 point scale.
●  * Example:
●  * row 2, col 3 (=12) says that the degradation at load idx 2
after
●  * 8 ticks is 12/128 (which is an approximation of exact
factor 3^8/4^8).
●  *
●  * With this power of 2 load factors, we can degrade the load
n times
●  * by looking at 1 bits in n and doing as many mult/shift
instead of
●  * n mult/shifts needed by the exact degradation.
●  */
● #define DEGRADE_SHIFT          7
● static const unsigned char
●     degrade_zero_ticks[CPU_LOAD_IDX_MAX] = {0, 8, 32, 64,
128};
● static const unsigned char
●     degrade_factor[CPU_LOAD_IDX_MAX][DEGRADE_SHIFT + 1] =
{
●         {0, 0, 0, 0, 0, 0, 0, 0},
●         {64, 32, 8, 0, 0, 0, 0, 0},
●         {96, 72, 40, 12, 1, 0, 0},
●         {112, 98, 75, 43, 15, 1, 0},
●         {120, 112, 98, 76, 45, 16, 2} };

```



```

• /*
•  * Update cpu_load for any missed ticks, due to tickless
•  idle. The backlog
•  * would be when CPU is idle and so we just decay the old
•  load without
•  * adding any new load.
•  */
• static unsigned long
• decay_load_missed(unsigned long load, unsigned long
missed_updates, int idx)
• {
•     int j = 0;
•
•     if (!missed_updates)
•         return load;
•
•     if (missed_updates >= degrade_zero_ticks[idx])
•         return 0;
•
•     if (idx == 1)
•         return load >> missed_updates;
•
•     while (missed_updates) {
•         if (missed_updates % 2)
•             load = (load * degrade_factor[idx][j]) >>
DEGRADE_SHIFT;
•
•         missed_updates >>= 1;
•         j++;
•     }
•     return load;
• }

```

- cpu_load[]含5条均线，反应不同时间窗口长度下的负载情况；主要供load_balance()在不同场景判断是否负载平衡的比较基准，常用为cpu_load[0]和cpu_load[1];
- cpu_load[index]对应的时间长度为{0, 8, 32, 64, 128}，单位为tick;
- 移动均线的目的在于平滑样本的抖动，确定趋势的变化方向;

系统级的负载计算calc_global_load_tick()

系统级的平均负载(load average)可以通过以下命令(uptime、top、cat /proc/loadavg)查看：

```

• mate20:/ # cat proc/loadavg && uptime
• 1.38 1.49 1.58 1/1085 20184
• 16:10:43 up 1 day, 2:29, 0 users, load average: 1.38, 1.49, 1.58

```

“load average:”后面的3个数字分别表示1分钟、5分钟、15分钟的load average。可以从几方面去解析load average：

- If the averages are 0.0, then your system is idle.

- If the 1 minute average is higher than the 5 or 15 minute averages, then load is increasing.
- If the 1 minute average is lower than the 5 or 15 minute averages, then load is decreasing.
- If they are higher than your CPU count, then you might have a performance problem (it depends).

最早的系统级平均负载(load average)只会统计runnable状态。但是linux后面觉得这种统计方式代表不了系统的真实负载；举一个例子：系统换一个低速硬盘后，他的 runnable负载还会小于高速硬盘时的值；linux认为睡眠状态

(TASK_INTERRUPTIBLE/TASK_UNINTERRUPTIBLE)也是系统的一种负载，系统得不到服务是因为io/外设的负载过重；系统级负载统计函数calc_global_load_tick()中会把(this_rq->nr_running+this_rq->nr_uninterruptible)都计入负载。

下面来看看具体的代码计算：

每个cpu每隔5s更新本cpu rq的(nr_running+nr_uninterruptible)任务数量到系统全局变量calc_load_tasks，calc_load_tasks是整系统多个cpu(nr_running+nr_uninterruptible)任务数量的总和，多cpu在访问calc_load_tasks变量时使用原子操作来互斥。

```

• /*
•  * Called from scheduler_tick() to periodically update this CPU's
•  * active count.
•  */
• void calc_global_load_tick(struct rq *this_rq)
• {
•     long delta;
•     /*判断5s更新周期是否到达*/
•     if (time_before(jiffies, this_rq->calc_load_update))
•         return;
•     /*计算本cpu的负载变化到全局变量calc_load_tasks中*/
•     delta = calc_load_fold_active(this_rq);
•     if (delta)
•         atomic_long_add(delta, &calc_load_tasks);
•     /*更新calc_load_update时间.LOAD_FREQ: (5*HZ+1), 5s*/
•     this_rq->calc_load_update += LOAD_FREQ;
• }

```

多个cpu更新calc_load_tasks，但是计算load只由一个cpu来完成，这个cpu就是tick_do_timer_cpu。在linux time一文中，我们看到这个cpu就是专门来更新时间戳timer的(update_wall_time())。实际上它在更新时间戳的同时也会调用do_timer()->calc_global_load()来计算系统负载。

核心算法calc_load()的思想也是：旧的load*老化系数 + 新load*系数

假设单位1为FIXED_1=2^11=2028，EXP_1=1884、EXP_5=2014、EXP_15=2037，load的计算：

$load = old_load * (EXP_? / FIXED_1) + new_load * (FIXED_1 - EXP_?) / FIXED_1$

```

• do_timer() -> calc_global_load()
• ↓
• void calc_global_load(unsigned long ticks)
• {
•     long active, delta;

```

```

•
•
•      /* (1) 计算的间隔时间为5s + 10tick,
•          加10tick的目的就是让所有cpu都更新完calc_load_tasks,
•          tick_do_timer_cpu再来计算
•      */
•      if (time_before(jiffies, calc_load_update + 10))
•          return;
•
•      /*
•       * Fold the 'old' idle-delta to include all NO_HZ cpus.
•       */
•      delta = calc_load_fold_idle();
•      if (delta)
•          atomic_long_add(delta, &calc_load_tasks);
•
•      /* (2) 读取全局统计变量 */
•      active = atomic_long_read(&calc_load_tasks);
•      active = active > 0 ? active * FIXED_1 : 0;
•
•      /* (3) 计算1分钟、5分钟、15分钟的负载 */
•      avenrun[0] = calc_load(avenrun[0], EXP_1, active);
•      avenrun[1] = calc_load(avenrun[1], EXP_5, active);
•      avenrun[2] = calc_load(avenrun[2], EXP_15, active);
•
•      calc_load_update += LOAD_FREQ;
•
•      /*
•       * In case we idled for multiple LOAD_FREQ intervals,
•       * catch up in bulk. */
•      calc_global_nohz();
•  }
•
•  |→
•
•  /*
•   *  $a_1 = a_0 * e + a * (1 - e)$ 
•   */
•  static unsigned long
•  calc_load(unsigned long load, unsigned long exp, unsigned
•  long active)
•  {
•      unsigned long newload;
•
•      newload = load * exp + active * (FIXED_1 - exp);
•      if (active >= load)
•          newload += FIXED_1-1;
•
•      return newload / FIXED_1;
•  }
•
•  #define FSHIFT      11      /* nr of bits of precision */

```

```

• #define FIXED_1      (1<<FSHIFT) /* 1.0 as fixed-point */
• #define LOAD_FREQ    (5*HZ+1)    /* 5 sec intervals */
• #define EXP_1        1884         /* 1/exp(5sec/1min) as
fixed-point */
• #define EXP_5        2014         /* 1/exp(5sec/5min) */
• #define EXP_15       2037         /* 1/exp(5sec/15min) */

```

对于cat /proc/loadavg的数值计算源码如下:

```

• #define LOAD_INT(x) ((x) >> FSHIFT)
• #define LOAD_FRAC(x) LOAD_INT((x) & (FIXED_1-1)) * 100)
•
• static int loadavg_proc_show(struct seq_file *m, void *v)
• {
•     unsigned long avnrun[3];
•
•     get_avenrun(avnrun, FIXED_1/200, 0);
•     /*其实还是直接获取系统全局变量,avnrun的数值在计算系统负载的时候已经计算了*/
•     seq_printf(m, "%lu.%02lu %lu.%02lu %lu.%02lu %ld/%d %d\n",
•         LOAD_INT(avnrun[0]), LOAD_FRAC(avnrun[0]),
•         LOAD_INT(avnrun[1]), LOAD_FRAC(avnrun[1]),
•         LOAD_INT(avnrun[2]), LOAD_FRAC(avnrun[2]),
•         nr_running(), nr_threads,
•         task_active_pid_ns(current)->last_pid);
•     return 0;
• }
•
• static int loadavg_proc_open(struct inode *inode, struct file *file)
• {
•     return single_open(file, loadavg_proc_show, NULL);
• }
•
• static const struct file_operations loadavg_proc_fops = {
•     .open      = loadavg_proc_open,
•     .read      = seq_read,
•     .llseek    = seq_lseek,
•     .release    = single_release,
• };
•
• static int __init proc_loadavg_init(void)
• {
•     proc_create("loadavg", 0, NULL, &loadavg_proc_fops);
•     return 0;
• }
• fs_initcall(proc_loadavg_init);

```

至此就计算完毕了.

